

OATSdb: vO

OATSdb: Getting Started

Avraham Leff

Yeshiva University

avraham.leff@yu.edu

COM3563: 2020

1. Introduction & “Nuts & Bolts”
2. **OATSdb** Base Classes & Interfaces
3. vo: Functional Requirements

Introduction & “Nuts & Bolts”



- ▶ **vo** will set up basic framework for **OATSdb** implementation!
- ▶ Nothing (certainly not software) is “cast in stone”, so you can always go back and fix things in subsequent milestones
 - ▶ But worth putting in the effort to start things off on a good foundation ☺
- ▶ Today's lecture provides both a general introduction to the **OATSdb** project and concepts and a discussion that's **specific to the vo milestone**

The requirements document for the **OATSdb** project is posted on Piazza: read it carefully!

- ▶ Consider this lecture an elaboration of some of the points discussed in the requirements document
 - ▶ Some students prefer a “slide format”: hence this lecture
 - ▶ More importantly, the lecture enables an interactive discussion!
- ▶ However: the requirements document provides motivation and context that I simply don't have space for in this lecture
 - ▶ The document should be considered the “master version” relative to the lecture
 - ▶ I'll update that document with useful feedback from you

- ▶ I've created a github repository named **OATSdb Base**
- ▶ It contains a set of base **classes & interfaces** for this project

- ▶ No excuses for **misspelling the APIs**
- ▶ Read the Javadoc **very, very carefully**: this is the “real” requirements document

- ▶ **All project milestones** will (re)implement the same set of interfaces
- ▶ Each milestone will provide an implementation with its own unique characteristics
- ▶ Repository code defines the `edu.yu.oatsdb.base` package
 - ▶ **vo**: you'll provide a `edu.yu.oatsdb.v0` package that implements these interfaces with “vo” semantics (today's lecture)

- ▶ VerifyLaptops open please: that you can download the repository
 - ▶ https://github.com/Yeshiva-University-CS/OATSdb_Base
- ▶ I'd strongly prefer that you **don't check-in this code** to your code base ...
- ▶ I **insist** that you **don't modify the code in any way!**
 - ▶ I will test your code only through the APIs defined in this package
 - ▶ Let's not take any chances breaking that 😊

See the requirements doc for the **required** directory layout and milestone **\$DIR** specification.

- ▶ A **Pragmatic Programmer** doesn't start a serious software project without a logging library
 - ▶ Logging via `println` just doesn't scale
- ▶ As usual, you can ignore this advice 😊
- ▶ That said: if you **do want to use a logging framework, you must use Apache Log4j 2**
 - ▶ Otherwise: I won't be able to add your non-JDK libraries to my classpath as I test your code
 - ▶ Result: will "break the build" ☹
 - ▶ With the usual implications ...

- ▶ If you **do use Log4j 2** in your code ...
- ▶ When **I execute your code**, your output will be generated and logged as well
 - ▶ This will aid debugging and help us understand what went wrong
 - ▶ Assuming, of course, that anything does go wrong ☺
- ▶ I'll bundle a *log4j2.properties* file in my build of your code, and link in the required libraries at runtime
- ▶ I've included a usable *log4j2.properties* file in the **OATSdb Base** repository
- ▶ Tweak as desired, but note this line:
`logger.file.name=edu.yu.oatsdb`
 - ▶ My code execution will only emit log output for packages nested under this package
- ▶ **Warning:** code that will be executed frequently, should be associated with a **DEBUG** level or below
 - ▶ You really, really, don't want to be emitting tons of output during "production"

Some miscellaneous observations

- ▶ I try to use JDK classes whenever possible, but find the `java.util.logging` APIs and output hard to use
- ▶ Every serious programmer writes at least two logging frameworks of their own before realizing that they should use a hardened, well-known, library instead ☺
 - ▶ You might as well “pick one, and stick with it”
 - ▶ `log4j2` is as good a choice as any
 - ▶ Let me know if you discover something better
- ▶ Configuration for **all logging libraries** can be a pain, ditto for understanding the framework's pieces
 - ▶ Suggestion: learn how to use at a **basic level** and don't try anything fancy
 - ▶ The Internet is your friend, but keep in mind that I'm using `log4j2`
 - ▶ **Stackoverflow is your special friend**

- ▶ All frameworks have multiple (severity) logging levels
 - ▶ Good frameworks make it easy to enable/disable logging at **package**, **class**, and **severity** granularities
- ▶ Separates logging API from logging **implementation**
 - ▶ You can plug-and-play different implementations (e.g., SL4J, *java.util.logging* etc)
- ▶ Even has support for different logging **APIs**
- ▶ **log4j2** (supposedly) has good performance, and low amount of “garbage collection” stress
- ▶ Advice: keep it simple
 - ▶ Pick anything that works and focus on **your code** 😊
- ▶ If you use maven, here is a sample “dependency” specification

```
<dependency>
  <groupId>org.apache.logging.log4j</groupId>
  <artifactId>log4j-core</artifactId>
  <version>2.8.2</version>
</dependency>
<dependency>
  <groupId>org.apache.logging.log4j</groupId>
  <artifactId>log4j-api</artifactId>
  <version>2.8.2</version>
</dependency>
```

- ▶ Here is the one exception to the “don’t modify your version of the **OATSdb** Base code”
- ▶ Because it’s so easy to make a mistake when using the factory methods of *OATSDbType* ...
 - ▶ (We’ll discuss these methods soon ...)
- ▶ ...I’ve added some Log4j2 code in that class
 - ▶ Result: the **OATSdb** Base code has a *compile* and *runtime* dependency on Log4j2
- ▶ If you don’t want to have a dependency on this framework, feel free to “comment out” the relevant lines of code
 - ▶ Be careful to not check this change into Git

- ▶ You don't have to write a single test 😊
- ▶ But: if you don't have a robust test suite, I can almost guarantee that you will fail 😞
- ▶ Key point: a **Pragmatic Programmer** doesn't start a serious software project without **including a test suite**
- ▶ For each milestone, I'll discuss some "test coverage" ideas that you may want to implement
 - ▶ That list will not be exhaustive!
 - ▶ You are responsible for translating requirements into tests

- ▶ Because of my automation scripts, your code can depend on only: the JDK, log4j2, and JUnit 4
- ▶ You are more than welcome to (**additionally**) set up a separate code base that uses other dependencies
 - ▶ Meaning: you can **develop** in an environment customized to your taste & needs
 - ▶ The code that you check into the “**OATSdb** tree” however: cannot rely on these extra dependencies
- ▶ Also: if you include test code in the repository code base, you are responsible for **not “breaking the build”**
- ▶ I also suggest that your tests **be in a different package** from the “package under test”
 - ▶ This ensures that you’re only testing **public** APIs
- ▶ Finally: keep your tests segregated by milestone: e.g., “vo” test package, “v1”
 - ▶ This will considerably simplify the complexity

You Have (Almost) Total Implementation Freedom!

This is all you have to provide 😊

```
public enum DBMSImpl implements DBMS {  
    Instance;  
  
    // Implement DBMS Interface  
}  
  
public enum TxMgrImpl implements TxMgr {  
    Instance;  
  
    // Implement TxMgr interface  
}
```

- ▶ Everything else should be “private” or “package-private”
- ▶ I will use the OATSDbType factory methods to instantiate your implementation

```
public static DBMS dbmsFactory(OATSDbType oatsdbType) {}  
public static TxMgr txMgrFactory(OATSDbType oatsdbType) {}
```

OATSdb Base Classes & Interfaces

- ▶ The base interface classes reference these Exception classes
- ▶ The Javadoc is self-explanatory: your task is to **throw the correct Exception at the appropriate time**
- ▶ The list:
 - ▶ `ClientNotInTxException.java`
 - ▶ `ClientTxRolledBackException.java`
 - ▶ `NotSupportedException.java`
 - ▶ `RollbackException.java`
 - ▶ `SystemException.java`

- ▶ **OATSDbType** declares an enum specifying values for supported **OATSdb** versions (milestones)
- ▶ Note the factory methods!
- ▶ You’ll use these methods to instantiate the *DBMS* and *TxMgr* **singleton Instance** implementations
 - ▶ And I will use these methods as I test your code
- ▶ The purpose of these factory methods is to **decouple the base class** from linkage to implementation code (such as vo)
 - ▶ Instead: “instantiate by name” on demand ☺
 - ▶ We’re using the magic of **Reflection**
 - ▶ Specifically: we’re using the Bloch “**enum as Singleton pattern**” idea)
 - ▶ Bloch’s **Effective Java** worth reading!

- ▶ The following interfaces will be used in **later milestones**
- ▶ **Ignore them** for vo
 - ▶ `ConfigurableDBMS.java`
 - ▶ `ConfigurablePersistentDBMS.java`

At any moment, an **OATSdb** client is or is not in a **transaction**

```
// is this code in a transaction?  
final String mapName = UUID.randomUUID().toString();  
final String key = "the question";  
final int value = 42;  
  
// what about this code?  
final Map<String, Integer> map1 =  
    dbms.createMap(mapName, String.class, Integer.class);  
map1.put(key, value);
```

- ▶ Vanilla Java code doesn't have to be (it's up to you) in a transaction
 - ▶ Example: the first code snippet
- ▶ Any code that invokes a *DBMS* API must be in a transaction
- ▶ And: all Map (container) methods must also be in a transaction
 - ▶ Example: the second code snippet

```
txMgr.begin();  
createAccounts(nAccounts); // application specific code  
txMgr.commit();
```

Transaction model:

- ▶ Client writes application code
- ▶ Wraps application code in a **transaction**
- ▶ How?
 - ▶ `txMgr.begin` demarcates the start of the transaction boundary
 - ▶ Either: `txMgr.commit` demarcates the end of the transaction boundary
 - ▶ Or: `txMgr.rollback` demarcates the end of the transaction boundary
- ▶ You had better **“close that transaction”** one way or the other
 - ▶ Or subsequent code will **continue to execute** in that transaction context

- ▶ Both methods terminate the transaction
 - ▶ Meaning: you'll have to start a **new transaction** afterwards if you want to wrap subsequent code in a transaction
- ▶ `commit` makes the work done by the application **visible** to other clients and transactions
 - ▶ Meaning: **persistent** beyond this transaction's work
- ▶ `rollback` “undoes” the work done by the application
 - ▶ Will **not be visible** to other clients and transactions
 - ▶ **Not even visible** to the client that rolled back the transaction!

- ▶ The **visibility** semantics are not for vo!
- ▶ vo only responsible for providing the appropriate **transaction boundary** semantics

- ▶ The **OATSdb** representation of a “transaction” is the *Tx interface*
 - ▶ This interface is minimal: only two methods, one returns the *TxStatus* and the other returns the *TxCompletionStatus*
- ▶ *TxStatus*: represents the possible “life-cycle states” that a transaction may be assigned
 - ▶ Very much a point in time statement: the next nano-second may move a transaction from e.g., ACTIVE to COMMITTING
- ▶ *TxCompletionStatus*: represents a more “stable” view of a transaction
 - ▶ Either it's COMMITTED, ROLLEDBACK, or NOT_COMPLETED
 - ▶ Under normal conditions, you'd expect that last state to change to one of the first two pretty quickly
- ▶ *TxCompletionStatus.java* includes static utility methods that map a *TxCompletionStatus* to a *TxStatus*

- ▶ Clients can only access the current transaction (if any) through the `TxMgr.getTx` method
- ▶ This fact, combined with the “minimality” of the `Tx` interface, implies that you have (almost) complete freedom in how you choose to implement a “transaction”
😊
 - ▶ Suggestion: start out “bare bones”, add internal function as needed
- ▶ Your `TxMgr` is responsible for driving a client's `Tx` through the appropriate `TxStatus` and `TxCompletionStatus` states
 - ▶ Based on the transaction “demarcation” methods we discussed previously
- ▶ As I test your code, I'll be verifying that your `Tx` implementation correctly reflects the state that's required to be associated with a transaction

- ▶ Consider the implications of a **test failure** ...
 - ▶ This is relevant to the “other guy”, not you 😊
- ▶ The transaction associated with your client thread is likely in an indeterminate state
 - ▶ At the least: your **internal data-structures** may be corrupted
- ▶ My test suite will therefore make an attempt to `rollback` your transaction
- ▶ This enables the **next test** to begin with a clean state
- ▶ Key point: your implementation of `rollback` (and `commit`) must do all the necessary data-structure cleanup!
 - ▶ Failure to do this properly (**even for vo**) will indeed result in cascading errors 😊
- ▶ Since I use `TxMgr.getTx` and `Tx.getStatus` to determine whether it's safe to issue a `rollback`, make sure that you provide a good implementation!

The Rules:

- ▶ A client thread can be associated with **at most one** transaction
- ▶ A client thread can only be associated with a transaction if it invoked `TxMgr.begin` and has not yet invoked `txMgr.commit` or `txMgr.rollback`
- ▶ Your **server implementation** must allow for **multiple client threads** to access its methods concurrently
- ▶ When a client invokes `TxMgr.begin`, your server is therefore responsible for creating a **unique association** between a given client thread and a `Tx` instance
- ▶ Your server is responsible for **removing that association** upon invocation of `TxMgr.commit` or `TxMgr.rollback`
- ▶ The implementation details of this process are up to you: but you must provide these semantics!

- ▶ Provides methods to “create” and “retrieve” named Maps
- ▶ Remember: (Generic) Maps are the **OATSdb** “table”
- ▶ MapEntries are the **OATSdb** “rows”
- ▶ **Q:** why do the Maps have to be **named**?
 - ▶ Vanilla Java maps are not named ...
- ▶ **A:** In vanilla Java code, you use **references** to distinguish between two Map instances
 - ▶ The purpose of **OATSdb** named Maps is to allow retrieval of previously persisted Maps
 - ▶ Where the original Map reference is **no longer available**
- ▶ Naming is **even more important** since your server must be able to distinguish between two Maps of the same type

- ▶ In your CS career, so far, you may have only used **generics** as a consumer
- ▶ Now you'll be **writing your own** generic classes!
- ▶ I urge you to get up to speed on this important topic
 - ▶ You need to understand the **semantics**, not just the syntax
 - ▶ **This** is a good tutorial
 - ▶ If it doesn't work for you, get a good textbook, use the Internet etc

```
Map<Long, Account> savings =  
    dbms.createMap(savingAccounts, Long.class, Account  
        .class);  
Map<String, Account> namedSavings =  
    dbms.getMap(namedSavingAccounts, String.class,  
        Account.class);
```

- ▶ Conceptually: these Maps have an existence that is independent of your program's (short-lived) execution
- ▶ Therefore
 - ▶ `createMap` will throw *IllegalArgumentException* if "name is already in use"
 - ▶ No subtlety here: exception thrown even if name is bound to a Map of a **different type**
 - ▶ `getMap` will throw *NoSuchElementException* if (named) Map does not exist

- ▶ These are conceptually **orthogonal** ideas
- ▶ **Database**: store the data and provide APIs to access that data
- ▶ **Transactions**: protect the DBMS data from corruption created by **concurrent multiple threads**
- ▶ The interaction between the *DBMS* and *TxMgr* occurs because **OATSdb** requires that all database data **must be accessed inside a transaction's scope**
 - ▶ This is how subsequent milestones will enable **concurrency** and **isolation** properties
- ▶ Hence the `ClientNotInTxException` declared on the DBMS methods
- ▶ Also: clients may only perform an operation on Map data when inside a transaction
 - ▶ The DBMS is responsible for throwing `ClientNotInTxException` if this requirement is violated

- ▶ There are well-defined nested transaction models
- ▶ **OATSdb** does not support any nested transaction model
 - ▶ Too much work ☺
 - ▶ Value (can be) dubious
- ▶ So: begin must throw a *NotSupportedException* if client begins a tx while already in a tx
- ▶ commit and rollback must throw *IllegalStateException* if client is not already in a tx

- ▶ The Maps returned by **DBMS** are `java.util.Map` instances
 - ▶ With a (possibly arbitrary) subsetting of the API
 - ▶ See below
- ▶ Use `put` to insert items
- ▶ Use `get` to retrieve items
- ▶ Use `remove` to remove items
- ▶ Conceptually: you'll be providing a “pass-through” implementation to a vanilla `java.util.Map`
 - ▶ **Not saying** that you must provide this implementation!

- ▶ The set **create, read, update and delete** database operations are affectionally known as **CRUD**
- ▶ These methods are the “Hello World” of database programming
- ▶ Many well-known “application builders” built their fame on making CRUD simple (especially for the web)
 - ▶ **Ruby On Rails**
 - ▶ Node.js + Express + MongoDB
 - ▶ No time now for snark ☺
- ▶ **Q:** where is the **update** method in the **OATSdb** API?

- ▶ In **OATSdb**, an “update” is done by changing a Map.Entry’s **value field**
 1. Acquire a reference to a Map.Entry through `get` or `put`
 2. Update the Map.Entry value field by changing its state
 3. Commit the transaction
- ▶ Key point: your main-memory **object reference** and the corresponding **database reference** are (conceptually) pointing to the same object!
 - ▶ Sometimes referred to as **Single-level-store semantics**
- ▶ Makes your life as an **OATSdb** provider trickier 😞
 - ▶ Because you must ensure that state changes made to any Map.Entry are visible to all clients that have acquired a reference to that Map.Entry

- ▶ You have to provide some magic 😊
- ▶ When a client invokes a Map method, **OATSdb** must know:
 - ▶ **Whether** the client is associated with a transaction?
 - ▶ **What transaction** the client is associated with?
 - ▶ **What objects** is the transaction associated with?
- ▶ Investigate: **Thread.current Thread**
- ▶ Investigate: **ThreadLocal**
- ▶ The former enables you to get your client's execution thread (so doesn't have to be supplied by the client explicitly)
- ▶ The latter enables you to group (arbitrary) data on a **per-Thread** basis

vo: Functional Requirements

vo provides:

- ▶ A functional “main-memory” database with supporting **CRUD** operations using `java.util.Map` API
 - ▶ Client can create **named** generic Maps
 - ▶ Client retrieve generic Maps by **name**
- ▶ No persistence: **all data disappears when your process finishes!**
- ▶ A transactional **interface**
 - ▶ You're not providing **atomicity** or **isolation**
 - ▶ In other words: **vo TxMgr** methods that **demarcate transaction boundaries** methods are “no-ops” from the perspective of ACID semantics
 - ▶ They don't change the state of the database in any way
- ▶ vo **does implement** transactional **boundaries**
- ▶ **OATSdb** vo **does enforce** the semantics that all DBMS methods must be accessed in a transaction
- ▶ **OATSdb** vo **does enforce** the semantics that all Map methods must be accessed in a transaction

You'll start implementing:

- ▶ DBMS: a “container of containers”
 - ▶ That is: must maintain all Maps created by the client
 - ▶ Identifying them by name
 - ▶ And: in a **generic** way
- ▶ TxMgr: map client threads to Tx instances
- ▶ You will become (more) familiar with Java generics
- ▶ You will become (more) familiar with Java Threads
- ▶ You will start building your test suites
- ▶ Note: **container concept**
- ▶ Client codes to “POJO” model, but container (**OATSdb**) is **transparently** mapping to more powerful/complicated implementation
 - ▶ Example: Map interface **knows nothing about transactions**, but your vo Map implementation requires that all clients invoking methods on it **be associated with a transaction**

You'll build on vo implementation for subsequent milestones

- ▶ You must make provision for client asking you to create (and store) an arbitrary generic Map
- ▶ You need a Map that associates a Map name with arbitrary generic – **not raw** – Map instances
- ▶ You need to store the “key class” and “value class” information
 - ▶ So you can validate that the client’s request can be met by the Map you’re storing for her
- ▶ Advice
 - ▶ Investigate generic **Wildcards**
 - ▶ Investigate “heterogeneous containers”
- ▶ Some useful pointers
 - ▶ Bloch, Effective Java (2nd edition) #29 “consider typesafe heterogeneous containers”
 - ▶ Bloch, Effective Java (2nd edition) #3: “Enforce the singleton property with a private constructor or an enum type”

- ▶ `java.util.Map` is an **interface**
- ▶ You may choose to write the **OATSdb** implementation class from scratch
 - ▶ Remember: you're returning a `Map` instance to your client
 - ▶ Must behave (at least for `CRUD` interface) like a `Map`
- ▶ But must be more than a `Map`!
 - ▶ At minimum: must be “transaction aware”
- ▶ Suggestion: investigate the **Proxy design pattern**
 - ▶ **Java support** for this pattern
- ▶ Lots of articles: e.g., **this one**

- ▶ You can & **should** use ideas in textbooks, Internet, videos
 - ▶ Be professional: document where you got your idea(s) from & how/whether you modified/extended
- ▶ You **absolutely may not** share ideas with fellow students
 - ▶ Code *a fortiori*
- ▶ There is **no one true way!**
 - ▶ I've done a "existence proof" that **OATSdb** can be implemented
 - ▶ You may well devise a better implementation

DBMS Happy Path

- ▶ Create a Map
- ▶ Retrieve a Map
- ▶ Suggest your test includes at least two Maps to make sure you're not cheating 😊
- ▶ put, then get Map.Entry instances

DBMS Edge Cases

- ▶ Create a Map **not in a transaction**
- ▶ Retrieve a Map **not in a transaction**
- ▶ CRUD **not in a transaction**
- ▶ What happens if you specify different key & value class for `getMap` than what you supplied for `createMap`?
- ▶ What does `createMap` do if that Map already exists?
- ▶ What does `getMap` do if that Map does not already exist?
- ▶ What if Map name is empty?

TxMgr

- ▶ Does **OATSdb** enforce “client must be in a transaction” for relevant TxMgr methods?
- ▶ When a client does a commit or rollback: is that thread **removed from the transaction** context?
- ▶ Does **OATSdb** enforce “no nested transactions”?

Note: Not much else to test for vo transactions ...

- ▶ vo **does not provide** ACID semantics
- ▶ Can you write code that proves that?
 - ▶ Easy: conceptually 😊
 - ▶ Hard: Computers are so fast that may not be trivial to write (efficient) test case that demonstrates “no transactions”

Suggestions For Getting Started (I)

- ▶ Even vo may seem initially daunting, so here are some suggestions (feel free to ignore 😊) that may help you get started
- ▶ Spend some time writing “throw-away” (“experimental”) code
 - ▶ Meaning: small chunks of code that are not integrated into an official package or class structure
- ▶ Some examples
 - ▶ Create a “generic container of containers”
 - ▶ Create a new generic Map and store it in that container
 - ▶ Retrieve a generic Map instance after it's been stored it in that container of containers
- ▶ Other examples
 - ▶ How do I access a client's thread?
 - ▶ Can I **associate state** with that thread for subsequent use?
 - ▶ Can I create a Map implementation class that behaves like a vanilla main-memory Map and yet allows me to do **something different** than a vanilla Map
 - ▶ Trivial function proves the point: print “hi mom” when get is invoked, print “hi dad” when put is invoked

Suggestions For Getting Started (II)

- ▶ In parallel with the previous activities
 - ▶ As you experiment with “getting a client’s thread”, create **two client threads** ...
 - ▶ Have each thread access the same piece of code in another class
 - ▶ Can the code in that other class **distinguish between those two threads**?
- ▶ Investigate the use of *java.util.concurrent.Future* interface (and related implementations) to create **multiple threads**
 - ▶ These “multiple threads” are your “multiple clients”
 - ▶ Have those threads access some other class
- ▶ Use “throw-away” code to quickly determine whether you’re on the right track!
- ▶ Read the Javadoc that I wrote for the **OATSdb** Base repository carefully!
- ▶ Read the references included in the lecture before you write any serious code: why reinvent the wheel?
 - ▶ Or at least figure out why you’re reinventing the wheel ☺

- ▶ Create a set of test method names (no implementation at first)
 - ▶ This enables you to express “what am I planning to build?”
 - ▶ We call this **test driven development**
- ▶ Once you have these “micro pieces” in place, you’ll have reasonable confidence that you can supply the vo **OATSdb** functional requirements
- ▶ Only afterward should you set up the **OATSdb** packages and start **integrating the “micro pieces”** into an implementation of the **OATSdb** interface

Introduction & “Nuts & Bolts”

OATSdb Base Classes & Interfaces

vo: Functional Requirements