

SQL: Finish Nested Queries & CUD Operation

COM 3563: Database Implementation

Avraham Leff


Yeshiva University

avraham.leff@yu.edu

COM3563: Fall 2020

Today's Lecture: Overview

1. More On SOME & ALL, Introduce EXISTS & NOT EXISTS
2. CUD Operations
3. Bonus Material: “Equivalent Queries”



- ▶ We're solidifying our grasp of SQL essentials ☺
- ▶ Previous lecture began the important topic of **nested queries**
 - ▶ Especially important because it gives the ability to develop "iteratively"!
- ▶ Today:
 - ▶ Finish the topic of nested queries
 - ▶ CUD operations

More On SOME & ALL, Introduce EXISTS & NOT EXISTS

Brain-Teasers: I

- ▶ The “truth value” of $x \text{ comparatorOp SOME } R \dots$
 - ▶ (R is a *relation*, *comparatorOp* is one of $<, \leq, >, =, \neq$)
- ▶ ...is equivalent to $\exists t \in R$ such that $(x \text{ comparatorOp } t)$

Spend a little time to make sure you agree with the following ...

```
1      (5 < some (0, 5, 6)) = true
2      (5 < some (0, 5)) = false
3      (5 = some (0, 5)) = true
4      (5 <> some (0, 5)) = true
```

- ▶ Note: $= \text{some}$ is equivalent to IN
- ▶ Note: $<> \text{some}$ is not equivalent to NOT IN

Brain-Teasers: II

- ▶ The “truth value” of x comparatorOp ALL R ...
 - ▶ (R is a *relation*, *comparatorOp* is one of $<, \leq, >, =, \neq$)
- ▶ ...is equivalent to $\forall t \in R$ (x comparatorOp t)

Spend a little time to make sure you agree with the following ...

```
1      (5 < all (0, 5, 6)) = false
2      (5 < all (6, 10)) = true
3      (5 = all (4, 5)) = false
4      (5 <> all (4, 6)) = true
```

- ▶ Note: $<>$ *all* is equivalent to NOT IN
- ▶ Note: $=$ *all* is not equivalent to IN

Correlated Nested Queries

- ▶ When a WHERE clause condition in a nested query references a relation attribute **declared in the outer query** ...
 - ▶ We say that the two queries are **correlated**
- ▶ Semantics: evaluate the nested query once for each tuple in the outer query
 - ▶ In contrast to the examples we've been looking at in which the nested query is evaluated **independently** of the outer query

```
1  -- For each EMPLOYEE tuple, evaluate the nested query, which retrieves the
2  -- Essn values for all DEPENDENT tuples with the same sex and first name
3  -- as that EMPLOYEE tuple ...
4  -- If the Ssn value of the EMPLOYEE tuple is in the result
5  -- of the nested query, then select that EMPLOYEE tuple
6  --
7  -- (Comment by me: "not a very useful query", pedagogic only)
8  SELECT E.Fname, E.Lname
9         FROM EMPLOYEE AS E
10        WHERE E.Ssn IN (SELECT D.Essn
11                        FROM DEPENDENT AS D
12                        WHERE E.Fname = D.Dependent_name
13                        AND E.Sex = D.Sex);
```

Introducing EXISTS

- ▶ EXISTS is a Boolean function and can therefore be used in a WHERE clause
- ▶ We use EXISTS to **check whether the result of a nested query is empty** or not
 - ▶ Return true *iff* the nested query result contains at least one tuple
- ▶ We use EXISTS (and NOT EXISTS) in conjunction with a **correlated nested query**
- ▶ Example below reimplements the previous query using EXISTS

```
1  SELECT E.Fname, E.Lname
2  FROM EMPLOYEE AS E
3  WHERE E.Ssn IN
4  (SELECT D.Essn
5   FROM DEPENDENT AS D
6   WHERE E.Fname = D.Dependent_name
7   AND E.Sex = D.Sex);
```

```
1  SELECT E.Fname, E.Lname
2  FROM EMPLOYEE AS E
3  WHERE EXISTS
4  (SELECT FROM WHERE *
5   DEPENDENT AS D
6   E.Ssn = D.Essn AND E.Sex = D.Sex
7   AND E.Fname = D.Dependent_name);
```


Empty Relation?

- ▶ **exists R** has the same truth value as $R \neq \emptyset$
- ▶ **not exists R** has the same truth value as $R = \emptyset$

Reimplement: *Find all courses taught in both the Fall 2009 semester and in the Spring 2010 semester*

```
1 SELECT course_id
2   FROM section AS S
3  WHERE semester = 'Fall' AND year = 2009
4     AND EXISTS (SELECT * FROM section AS T
5                  WHERE semester = 'Spring' AND year=2010
6                  AND S.course_id = T.course_id);
```

- ▶ **Correlation name:** variable S in the outer query
- ▶ **Correlated sub-query:** inner sub-query that uses a “correlation name” from an outer query

Using NOT EXISTS to Test For Set Containment (“Superset”)

Observation: given sets X and Y

- ▶ $X - Y = \emptyset \iff X \subseteq Y$
- ▶ In other words: relation $Y \supseteq X$

Implication: *Is set X a superset of Y ? ...*

Can be implemented in SQL as:

```
1      not exists (X except Y)
```

Set Containment: Example

Find all students who have taken all courses offered in the Biology department

```
1  -- Examine each student: does the set of all Biology courses
2  -- taken by this student contain the set of all Biology
3  -- courses?
4  SELECT DISTINCT S.ID, S.name
5  FROM student AS S
6  WHERE NOT EXISTS
7      -- all courses offered in Biology
8      ((SELECT course_id
9         FROM course
10        WHERE dept_name = 'Biology')
11     EXCEPT
12     -- all courses taken by a particular student
13     (SELECT T.course_id
14      FROM takes AS T
15      WHERE S.ID = T.ID));
```

Which customers have ordered something?

Customer

<u>id</u>	name
1	Joe
4	Mary
3	Scott
6	Elizabeth

Order

<u>id</u>	customer_id	item
105	4	Shoes
107	4	Pants
108	1	Pants
109	3	Tie

$\{c \mid \text{Customer}(c) \wedge$
 $(\exists o. \text{Order}(o) \wedge$
 $c.\text{id} = o.\text{customer_id})\}$



<u>id</u>	name
1	Joe
4	Mary
3	Scott

SELECT id, name
FROM Customer c
WHERE EXISTS (SELECT *
FROM Order o
WHERE c.id = o.customer_id);

EXISTS (subquery)
returns a Boolean.

Subquery's SELECT
irrelevant.

Which customers have not ordered anything?

Customer

<u>id</u>	name
1	Joe
4	Mary
3	Scott
6	Elizabeth

Order

<u>id</u>	customer_id	item
105	4	Shoes
107	4	Pants
108	1	Pants
109	3	Tie

$$\{c \mid \text{Customer}(c) \wedge (\neg \exists o. \text{Order}(o) \wedge c.\text{id} = o.\text{customer_id})\}$$


<u>id</u>	name
6	Elizabeth

```
SELECT id, name
FROM Customer c
WHERE NOT EXISTS (SELECT *
                  FROM Order o
                  WHERE c.id = o.customer_id);
```

Efficiency

```
SELECT id, name  
FROM Customer c  
WHERE EXISTS (SELECT *  
               FROM Order o  
               WHERE c.id = o.customer_id);
```

At first satisfying Order,
stop and return TRUE.

```
SELECT id, name  
FROM Customer c  
WHERE NOT EXISTS (SELECT *  
                  FROM Order o  
                  WHERE c.id = o.customer_id);
```

At first satisfying Order,
stop and return FALSE.

Logical equivalence of quantifiers

$$\neg \exists x . R(x) = \forall x . \neg R(x)$$

$$\neg \forall x . R(x) = \exists x . \neg R(x)$$

Universal quantification – no FORALL()!

Generally a less useful idea in SQL.

- $\exists x. R(x)$ – Exists some Order o such that the Customer c ordered it.
- $\forall x. R(x)$ – For all Orders o, Customer c ordered it.

$$\forall x. R(x) = \neg \exists x. \neg R(x)$$

```
SELECT id, name
FROM Customer c
WHERE FORALL (SELECT *
              FROM Order o
              WHERE c.id = o.customer_id);
```



```
SELECT id, name
FROM Customer c
WHERE NOT EXISTS (SELECT *
                  FROM Order o
                  WHERE c.id <> o.customer_id);
```


Another form of existential quantification

```
SELECT id, name  
FROM Customer c  
WHERE EXISTS (SELECT *  
               FROM Order o  
               WHERE c.id = o.customer_id);
```

```
SELECT id, name  
FROM Customer  
WHERE id = ANY (SELECT customer_id  
                FROM Order);
```

The operator is
= ANY()

Another form of universal quantification

```
SELECT *  
FROM Product p1  
WHERE p1.price >= ALL (SELECT p2.price  
                        FROM Product p2);
```



```
SELECT *  
FROM Product p1  
WHERE NOT (p1.price < ANY (SELECT p2.price  
                           FROM Product p2));
```

Nested Queries: Can Be Used With FROM Clause

- ▶ We've focused our **nested query** discussion on usage with **WHERE** clause
- ▶ But: can also be used with a query's **FROM** clause
- ▶ After all: **SELECT-FROM-where** **returns a relation**
 - ▶ So: we can plug the relation returned by a nested query into any place where an SQL query expects to see a relation
 - ▶ Including the **input to a FROM clause** 😊
- ▶ Example: *"Find the average instructor salary of those departments where the average salary is greater than \$42,000."*

```
1 SELECT dept_name, avg_salary
2   FROM (SELECT dept_name, AVG (salary) AS avg_salary
3         FROM instructor
4         GROUP BY dept_name)
5  WHERE avg_salary > 42000;
```

Naming Nested Query Artifacts

- ▶ You can name a nested query's result set
- ▶ You can rename the attributes from a nested query's result set

```
1 SELECT dept_name, avg_salary
2     FROM (SELECT dept_name, avg (salary)
3            FROM instructor
4            GROUP by dept_name)
5         AS dept_avg (dept_name, avg_salary)
6     WHERE avg_salary > 42000;
```

Nasty PostgreSQL Bug

- ▶ Q: what's wrong with this query?

```
1  SELECT COUNT (made_only_recharge) AS made_only_recharge
2  FROM (
3      SELECT DISTINCT (identifiant) AS made_only_recharge
4      FROM cdr_data
5      WHERE CALLEDNUMBER = '0130'
6      EXCEPT
7      SELECT DISTINCT (identifiant) AS made_only_recharge
8      FROM cdr_data
9      WHERE CALLEDNUMBER != '0130'
10 )
```

- ▶ Per [this discussion](#), the query is OK on Oracle, breaks for PostgreSQL ☹
- ▶ A: there's nothing wrong with the query **except that PostgreSQL insists that each sub-query relation in a FROM clause must be given a name**
 - ▶ Even if the name is never referenced! ☹
- ▶ But the truth is: good SQL style suggests that you always “alias” your sub-query
 - ▶ Never hurts to make your code more clear ☺

Nested Queries: Can Be Used With SELECT Clause

- ▶ You can even use a nested query in the SELECT clause!
- ▶ We refer to this scenario as a **scalar sub-query**
 - ▶ “Scalar”: because you expect only a **single value** as the result
 - ▶ “Scalar”: because **runtime error** if sub-query returns more than one result tuple 😊
- ▶ Example: *“List all departments along with the number of instructors in each department”*

```
1  SELECT dept_name, (SELECT count(*)
2                        FROM instructor
3                        WHERE department.dept_name = instructor.
                           dept_name)
4                        AS num_instructors
5  FROM department;
```


Introduction

- ▶ We've been polishing our SQL **query** (the R from CRUD) skills ...
- ▶ But ignored CUD operations 😞
- ▶ One reason: CUD operations are an SQL “add-on” to the theory of relational databases
 - ▶ A necessary feature for “commercial” systems
 - ▶ But not fundamentally necessary 😊
- ▶ Another reason: sophisticated CUD operations incorporate queries to specify the relational context to which the operation will be applied
- ▶ We're now ready to tackle this topic

Create

- ▶ Simplest form: use to add one or more tuples to a relation
- ▶ As we've discussed:
 - ▶ Constraints on data types are observed automatically
 - ▶ Integrity constraints (part of the DDL specification) are enforced
- ▶ Lazy version (don't do this!):
 - ▶ Can specify the attribute names **implicitly** by supplying the attribute values in **the same order** as they're specified in the CREATE TABLE DDL

```
1 insert into instructor
2     values ('10211', 'Smith', 'Biology', 66000)
```

- ▶ Better:
 - ▶ Specify attribute names **explicitly**, such that they are associated one-to-one with the supplied values

```
1 insert into instructor(name, ID, dep_name, salary)
2     values ('Smith', '10211', 'Biology', 66000)
```

Create With Missing Values

- ▶ If you omit values of certain columns
 - ▶ Default values will be used (per prior DDL statement)
 - ▶ Or the NULL value will be used

```
1 insert into instructor(name, ID, dep_name)
2 values ('Smith', '10211', 'Biology')
```

- ▶ The above example (implicitly) **does not supply** a value for the *salary* attribute
- ▶ As with previous slide: better to make explicitly state that you're supplying NULL for the *salary* attribute

Inserting From Another Relation

- ▶ Example of the beauty that comes from a set-based programming model
- ▶ Input to the “insert” operation is defined in terms of **tuples**
- ▶ Queries **return tuples**
- ▶ So: we can insert the result of a query!

```
1  INSERT into student
2    SELECT ID, name, dept_name, 0
3    FROM instructor
```

- ▶ This SQL add all instructors to the *student* relation, and sets their *tot_creds* attribute set to 0
- ▶ Can insert any result of an query, as long as compatible, into a table

Source of Insert Operation Can Be Any Query

- ▶ You can use the full power of a query to drive an “insert” include the SELECT FROM WHERE version
- ▶ Example: *“Make each student in the Music department who has earned more than 144 credit hours an instructor in the Music department with a salary of \$18,000”*

```
1  INSERT INTO instructor
2      SELECT ID, name, dept_name, 18000
3      FROM student WHERE dept_name = 'Music' AND
         total_cred > 144;
```

- ▶ Semantics: DBMS **fully evaluates** the query results **before** inserting results into relation
- ▶ Otherwise the following statement could cause an **infinite loop ...**

```
1      insert into table1 select * from table1
```

Adding records to a table

```
INSERT INTO Product
VALUES
  ('MiniGizmo', 15.99, 'GizmoWorks'),
  ('MiniWidget', 21.99, 'WidgetsRUs');

INSERT INTO Product (prod_name, prod_manufacturer)
VALUES
  ('NanoWidget', 'WidgetsRUs');
```

Explicitly listing column names is safer. Protects against table column changes.

Product

prod_name	prod_price	prod_manufacturer
'Gizmo'	\$19.99	'GizmoWorks'
'Powergizmo'	\$39.99	'GizmoWorks'
'Widget'	\$19.99	'WidgetsRUs'
'HyperWidget'	\$203.99	'Hyper'



Product

prod_name	prod_price	prod_manufacturer
'Gizmo'	\$19.99	'GizmoWorks'
'Powergizmo'	\$39.99	'GizmoWorks'
'Widget'	\$19.99	'WidgetsRUs'
'HyperWidget'	\$203.99	'Hyper'
'MiniGizmo'	\$15.99	'GizmoWorks'
'MiniWidget'	\$21.99	'WidgetsRUs'
'NanoWidget'	NULL	'WidgetsRUs'

Update

```
1      update instructor set salary = salary * 1.03
2      where salary > 100000
3
4      update instructor set salary = salary * 1.05
5      where salary <= 100000
```

- ▶ Use the UPDATE command to modify attribute values of one or more tuples
- ▶ Use a WHERE-clause to specify the tuples that will be modified
 - ▶ In other words: you use a **query** to specify the set of tuples that will be updated
- ▶ Then use a SET-clause to specify the attributes to be modified and **their new values**
 - ▶ Can selectively update a **subset of attributes**
- ▶ As discussed before, DBMS will enforce the referential and other integrity constraints that were specified as part of the DDL

Queries Are Completely Integrated With Updates

```
1 UPDATE instructor
2   SET salary = salary * 1.05
3   WHERE salary < (select AVG (salary) FROM instructor);
```

- ▶ Just because we're doing **declarative programming**, don't forget your basic programming skills!
- ▶ Example: **statement order** matters!
- ▶ Consider this update: *Increase salaries of instructors whose salary is over \$100,000 by 3%, and all others by a 5%*

```
1 UPDATE instructor
2   SET salary = salary * 1.05 WHERE salary <= 100000;
3 UPDATE instructor
4   SET salary = salary * 1.03 WHERE salary > 100000;
```

Q: can you spot the bug?

Statement Order Matters

```
1 UPDATE instructor
2   SET salary = salary * 1.05 WHERE salary <= 100000;
3 UPDATE instructor
4   SET salary = salary * 1.03 WHERE salary > 100000;
```

- ▶ A: some instructors may get two raises 😊
- ▶ Rather than rearranging the statement order, better to use SQL's **conditional update** syntax: the CASE statement

```
1 UPDATE instructor
2   SET salary = CASE
3       WHEN salary <= 100000 THEN salary *
4           1.05
5       ELSE salary * 1.03
6   END
```


Delete

- ▶ Use the DELETE command to remove tuples from a relation
- ▶ Include a WHERE-clause to specify the tuples to be deleted
- ▶ As discussed previously, the DBMS will enforce referential integrity
 - ▶ Example: quite difficult to delete tuples from a randomly selected table in the university database
- ▶ As with updates, a single statement can only delete tuples from a single table
- ▶ As with updates, if you specify CASCADE in the referential integrity constraint, a single statement can effect multiple tables 😊

Delete

- ▶ Set theory strikes again!
- ▶ If you **omit the WHERE-clause** from a DELETE statement, you're implicitly specifying "delete all tuples in the relation"
 - ▶ Equivalent to: WHERE TRUE, because every tuple satisfies the "empty condition"
 - ▶ Result: the table becomes an **empty table**
 - ▶ So: "be careful when you type" 😊
- ▶ *Delete all instructor tuples*

```
1 delete from instructor
```

- ▶ In general: the number of deleted tuples depends on the number of tuples in the relation that satisfy the WHERE-clause

```
1 delete from instructor where dept_name= 'Finance'
```

Adding data to a table from a file

Server file: `COPY Product FROM '/my/path/product.csv' CSV HEADER;`

Must be logged in as a superuser, e.g., postgres. Server app must have file permission.

Client file: `\copy Product FROM '/my/path/product.csv' CSV HEADER`

Must use psql.

<u>prod_name</u>	prod_price	prod_manufacturer
prod_name, prod_price, prod_manufacturer		
Gizmo,19.99,GizmoWorks		
Powergizmo,39.99,GizmoWorks		
Widget,19.99,WidgetsRUs		
HyperWidget,203.99,Hyper		



<u>prod_name</u>	prod_price	prod_manufacturer
Gizmo	\$19.99	GizmoWorks
Powergizmo	\$39.99	GizmoWorks
Widget	\$19.99	WidgetsRUs
HyperWidget	\$203.99	Hyper

Deleting data

```
DELETE FROM Student;
```

```
DELETE FROM Student  
WHERE ...;
```

However – often **never** want to delete!

- E.g., former customers, discontinued products

Soft delete:

- Have column to mark data as “inactive”.

Hard delete:

- Get rid of it!

Create, update, delete

Tables

```
CREATE TABLE Student (...);
```

```
ALTER TABLE Student  
RENAME TO Scholar;
```

```
DROP TABLE Student;
```

Columns

```
ALTER TABLE Student  
ADD COLUMN zip INT;
```

```
ALTER TABLE Student  
RENAME COLUMN zip TO zip_code;
```

```
ALTER TABLE Student  
ALTER COLUMN zip VARCHAR(9);
```

```
ALTER TABLE Student  
DROP COLUMN state;
```

Constraints

Similar. Later.

Changing schemas affects other code

Obvious: Is a table or column is defined?

More subtle:

- `SELECT *`
- `INSERT INTO` – without listing column names

How can you know what's defined?

```
SELECT * FROM Table;
```

psql command	What it does
<code>\d</code>	List defined items. (So far, that's just tables.)
<code>\dt</code>	List defined tables.
<code>\d <i>tablename</i></code>	Give details about table definition.

How to change schemas?



- Create new tables.
- Create new queries using the new tables.
- Copy or compute data for the new tables.
- Switch to using new tables and queries.

} Can take substantial time.
Lots of data. Possibly
distributed & replicated.

What about data that arrives during this process?

Traditional – shutdown

Annoying!



Best – invisible to user

A simple version:

- Create *shadow table* that includes schema changes.
- Create *triggers* in original table that forwards data updates to shadow table.
- Copy original table's data to shadow table.
- Rename shadow table to replace original table.

Can choose to only let some users see the new version (A/B testing).

Some DBMSes provide support for such *online* schema changes.

Bonus Material: “Equivalent Queries”

```
SELECT DISTINCT c.id, name
FROM Customer c
INNER JOIN Order o ON c.id = customer_id;
```

```
SELECT id, name
FROM Customer
WHERE id IN (SELECT customer_id
             FROM Order);
```

```
SELECT id, name
FROM Customer c
WHERE (SELECT Count(*)
       FROM Order o
       WHERE c.id = o.customer_id)
>= 1;
```

```
SELECT id, name
FROM Customer c
WHERE EXISTS (SELECT *
              FROM Order o
              WHERE c.id = o.customer_id);
```

```
SELECT id, name
FROM Customer
WHERE id = ANY (SELECT customer_id
                FROM Order);
```

```
SELECT id
FROM Customer
WHERE EXISTS (SELECT *
              FROM Order
              WHERE Customer.id = customer_id);
```

```
SELECT id
FROM Customer
INTERSECT
SELECT customer_id
FROM Order;
```

Today's Lecture: Wrapping it Up

More On SOME & ALL, Introduce EXISTS & NOT EXISTS

CUD Operations

Bonus Material: “Equivalent Queries”

- ▶ Textbook discusses **nested queries** in Chapter 3.8
- ▶ Textbook discussed **CUD operations** in Chapter 3.9

Do an end-to-end reading of **Chapter 3**