

# From E-R Model to Relational Database

COM 3563: Database Implementation

Avraham Leff

Yeshiva University

*avraham.leff@yu.edu*

COM3563: Fall 2020

# Today's Lecture: Overview

1. From E-R to Relational
2. Schema Redundancy
3. Review Of Common Relationship Set  $E-R \Rightarrow RDB$  Mappings

## Previous Lecture: The E-R Model

- ▶ E-R is valuable because it's a **semi-formal design technique based upon informal semantics** 😊
- ▶ When using E-R, our goal is to create schemas that accurately represent and formalize the information we acquired during the “requirements gathering” phase
  - ▶ The E-R modeling is valuable even though we know that complete accuracy isn't always possible
  - ▶ Tip-off: lots of “natural language” annotations
  - ▶ Tip-off: specialized E-R diagrams that are supported only by a given modeling tool
- ▶ The E-R modeling process
  1. Identify **entity sets** and their attributes
  2. Identify **relationships between entity sets**, and their attributes
  3. Identify maximum and minimum **relationship cardinality**
  4. Identify **weak entity sets**

# Today's Plan



- ▶ The real benefit of E-R modeling stems from the ability to generate “real” RDB schema from the E-R model
  - ▶ Implication: we can iteratively refine the “high-level” E-R model, based on the concrete RDB schema that it generates
- ▶ The two data-models are sufficiently similar that we can “compile down” the higher-level model (E-R) to a lower-level (RDB) “executable” version
- ▶ Key point: this transformation can be applied in quasi-algorithmic fashion



## Basic Idea

- ▶ E-R in a nutshell: **entity sets** and **relationship sets** ☺
- ▶ Entity sets and relationship sets can both be transformed **cleanly** to RDB schemas
  - ▶ E-R diagram → collection of RDB schemas
  - ▶ With minimal “hackery” ☺
- ▶ For each entity set and relationship set **there is a unique schema** that is assigned the name of the corresponding entity set or relationship set
- ▶ Each schema has a number of columns (generally corresponding to **attributes**), which have unique names
- ▶ The non-intuitive part is that we’re **converting relationship sets to schemas** (“tables”)!


# Representing Entity Sets

- ▶ Easy: a **strong entity set** reduces to a schema with the same attributes

```
1      student(ID, name, tot_cred)
```

- ▶ Less obvious: a **weak entity set** becomes a table that includes a column for the primary key of the identifying strong entity set
  - ▶ Example from last lecture: concept of a **section** doesn't have independent existence apart from a **course**
  - ▶ Discriminator for **section**: (*section\_id*, *semester*, *year*)
  - ▶ Primary key of **course** is *course\_id*

```
1  section (course_id, section_id, semester, year)
```



## Weak Entity Sets


- ▶ **Q:** How does a relational database capture the concept that a “weak entity set” is **existence dependent** on some “strong entity set”?
- ▶ **A:** By creating a **foreign-key constraint** on the corresponding “weak” relation specifying that
  - ▶ Each of its attributes “derived from” the primary key of the “strong” relation **reference** the primary key of that relation
  - ▶ Example: *course\_id* in schema *section* has a foreign key constraint on the primary key of the *course* schema
  - ▶ Can also add an **integrity constraint** on this *section* foreign-key constraint specifying an ON DELETE CASCADE **option**
    - ▶ This ensures that deleting a *course* tuple results in **all associated** *section* tuples being deleted as well



## E-R $\Rightarrow$ RDB Mapping (Steps 1 & 2)

- ▶ Step 1: mapping of strong entity sets
  - ▶ For each entity set  $E$  in the E-R schema, create a relation  $R$  that includes all the simple attributes of  $E$
  - ▶ Choose one of the key attributes of  $E$  as the primary key for  $R$
  - ▶ If the chosen key of  $E$  is **composite**, the set of simple attributes will form the primary key of  $R$
- ▶ Step 2: mapping of weak entity sets
  - ▶ For each weak entity set  $W$  in the E-R schema with owner entity set  $E$ , create a relation  $R$  and include all simple attributes of  $W$  as attributes of  $R$
  - ▶ Also, include as foreign key attributes of  $R$  the primary key attribute(s) of the relation(s) that correspond to the owner entity set(s)
  - ▶ The primary key of  $R$  is the combination of the primary key(s) of the owner(s) and the partial key of the weak entity set  $W$ , if any

# Entity Sets With Composite Attributes



- ▶ **Composite attributes** are “flattened” by creating a separate attribute for each component attribute
- ▶ Example: entity set ***instructor*** with composite attribute ***name*** composed of component attributes ***first\_name*** and ***last\_name*** → **two schema attributes**
  - ▶ ***name\_first\_name*** and ***name\_last\_name***
- ▶ Can omit the “name prefix” if no resulting ambiguity
  - ▶ OK to transform ***name\_first\_name*** into ***first\_name***
- ▶ (**Temporarily** ignoring) the issue of **multivalued** attributes, extended ***instructor*** schema is shown below

```
1 instructor(ID, first_name, middle_initial, last_name, street_number,  
2 street_name, apt_number, city, state, zip_code, date_of_birth)
```

## MultiValued Attributes

- ▶ Example from last lecture: **phone numbers** (because *instructor* can have 0, 1, or more phone numbers)
- ▶ A **multivalued** attribute  $M$  of an entity  $E$  is represented by a separate schema  $EM$
- ▶ Schema  $EM$  has primary key attributes corresponding to the **primary key** of  $E$  and an **attribute** corresponding to multivalued attribute  $M$
- ▶ Example: multivalued attribute `phone_number` of `instructor` is represented by


```
1      instructor_phone = (ID, phone_number)
```

- ▶ **Each value** of the multivalued attribute maps to a **separate tuple** of the relation on schema  $EM$
- ▶ Example: `instructor` entity with primary key 22222 and phone numbers 456-7890 and 123-4567 maps to **two tuples**
  - ▶ (22222, 456-7890)
  - ▶ (22222, 123-4567)
- ▶ Much more fun to work with NOSQL databases for multivalued attributes 😊

## Schema Generated From Multivalued Attributes

- ▶ Transformation process created a new schema from the multivalued attribute *phone number*
- ▶ We seem to have **lost the information** that *instructor\_phone* was derived from the *instructor* entity set
- ▶ **Q:** can you suggest how we might somehow **keep this information**?
- ▶ **A:** Create a **foreign-key constraint** on the new schema
  - ▶ Attribute generated from primary key of entity set references the schema generated from that entity set
  - ▶ For *instructor\_phone* schema, place foreign-key constraint such that attribute *ID* references the *instructor* relation

# Representing Relationship Sets



- ▶ A relationship set (independently of the participating entity sets) is transformed into a schema
- ▶ **Many-to-many** relationship set is represented as a schema with
  - ▶ Attributes for the **primary keys of the two participating entity sets** and
  - ▶ All descriptive attributes of the relationship set itself
- ▶ Example from last lecture: relationship set *advisor* is transformed into

```
1      advisor = (s_id, i_id)
```

## E-R $\Rightarrow$ RDB Mapping (Steps 3 & 4)

- ▶ Step 3: mapping of multi-valued attributes
  - ▶ For each multi-valued attribute  $A$ , create a new relation  $R$
  - ▶ This relation  $R$  will include an attribute corresponding to  $A$ , plus the primary key attribute  $K$  (as a foreign key in  $R$ ) of the entity set that has  $A$  as an attribute
  - ▶ The primary key of  $R$  is the combination of  $A$  and  $K$ 
    - ▶ If the multi-valued attribute is composite, include its simple components in the primary key
- ▶ Step 4: mapping of binary  $M : N$  relationship types
  - ▶ For each binary  $M:N$  relationship set  $R$ , create a new (“relationship relation”) relation  $S$  to represent  $R$
  - ▶ The primary keys of the participating entity sets become foreign key attributes in  $S$ 
    - ▶ The combination of these primary keys form the primary key of  $S$
  - ▶ Include any attributes of the  $M : N$  relationship set itself as attributes of  $S$

## E-R $\Rightarrow$ RDB Mapping (Step 5): Mapping Binary 1 : N Relationships

- ▶ For each binary 1 : N relationship type  $R$ , create a relation  $S$  to represent the participating entity set at the “N-side” of the relation
  - ▶ Example: given an *Employee*  $\Rightarrow$  *Department* WORKS\_IN relationship ...
  - ▶ ...*Employee* is the “N-side” of the relation
- ▶ Include the primary key of the relation  $T$  that represents the **other entity set** participating in  $R$  as a foreign key in  $S$ 
  - ▶ In our example: include “department\_number” as a **foreign key** in *Employee*
- ▶ Include any simple attributes of the 1 : N relation type as attributes of  $S$

## E-R $\Rightarrow$ RDB Mapping (Step 6): Mapping Binary 1 : 1 Relationships (I)

- ▶ For each binary 1 : 1 relationship type  $R$  in the E-R schema, identify the relations  $S$  and  $T$  that correspond to the entity types participating in  $R$
- ▶ There are three possible approaches:
  - ▶ **Foreign key (“2 relations”)** approach: choose one of the relations (e.g.,  $S$ ) and include the primary key of  $T$  as a foreign key in  $S$ 
    - ▶ Ideally: the entity type playing the role of  $S$  should have a **total participation** semantics in  $R$
  - ▶ **Merged relation (“1 relation”)** approach: merging  $S$  and  $T$  into a new **single relation**
    - ▶ This approach is reasonable when both participations are total
    - ▶ The relationship information is **now modeled implicitly**



- ▶ A third approach is possible as well
  - ▶ **Cross-reference** (“3 relations”) approach: model the relationship explicitly by creating a third relation  $X$
  - ▶  $X$  “cross-references” the primary keys of  $S$  and  $T$ 
    - ▶ This is sometimes called the “relationship relation” approach
    - ▶ The attributes of  $S$  and  $T$  are mapped directly from the participating entity types: the relationship between them is captured as “primary key” duples in  $X$

## Relationship Sets & Primary Keys: Summary (I)

- ▶ The naive approach for determining the **primary key** of the schema that implements the E-R relationship is to use the **union of the participating entity sets**
  - ▶ Use this as a “first approximation” approach for **non-binary** relationship
- ▶ But we can do better than “union of primary keys” by taking advantage of the fact that (many) binary relationships express a **partial** or **total** function between the two (participating) entity sets
- ▶ **Binary One-to-One**: primary key of **either** entity set becomes the primary key of the relation
  - ▶ Do you see why it doesn't matter which one we choose?
- ▶ **Binary Many-to-One (or One-to-Many)**: primary key of the entity set on the “many” side of the relationship becomes the primary key of  $R'$


## Relationship Sets & Primary Keys: Summary (II)

- ▶ **Binary Many-to-Many:** **union** of the primary keys of the participating entity sets becomes the primary key of  $R'$ 
  - ▶ This is the “relationship relation” approach that we mentioned as a possibility for 1 : 1 binary relationships
  - ▶ For  $M : N$  relationships, this is the only approach possible
- ▶ **Note:** the above rules for binary relationships simply express the fundamental semantics of “primary keys”
  - ▶ A primary key **uniquely** identifies a tuple
  - ▶ By understanding “partial” and “total” participation semantics, we optimize the amount of **primary key information** that’s needed to encode the relationship between the participating entity sets

## Relationship Sets: Summary of Foreign-Key Constraints Semantics


- ▶ Key point: these rules should be intuitive
  - ▶ A **foreign-key** is a column or group of columns that serve as a “**cross-reference**” or “**link**” between database tables
  - ▶ The “cross-reference” is implemented by specifying the **primary key** of another table
- ▶ Assume that each participating entity set  $E_i$  has a corresponding relational schema  $E'_i$
- ▶ Recall:  $R'$  is the **relation schema derived from relationship set  $R$**
- ▶ For every set of attributes in schema  $R'$  derived from primary key attributes of  $E'_i$ , **create a foreign-key constraint from  $R'$**  such that
  - ▶ The attributes of  $R'$  derived from the primary key attributes of  $E'_i$  ...
  - ▶ ...Are declared to reference the primary key of the relation schema  $E'_i$

## Relationship Set: Example (I)



- ▶ Relationship set *advisor* associates
  - ▶ *instructor* entity set with primary key *id*
  - ▶ *student* entity set with primary key *id*
  - ▶ In a one-to-many (instructor-to-student) binary relationship
- ▶ Relationship has no attributes of its own so *advisor* schema only has two attributes
  - ▶ Union of  $E_i$  **primary key** attributes
  - ▶ Must rename them as *instructor\_id* and *student\_id*
  - ▶ Primary key of *advisor* is the primary key of the “many side” of the relationship
    - ▶ Or: *student\_id*

## Relationship Set: Example (II)



- ▶ Create **two** foreign-key constraints in *advisor* schema
  - ▶ Attribute *instructor\_id* references primary key of *instructor*
  - ▶ Attribute *student\_id* references primary key of *student*




# Schema Redundancy



# Introduction

- ▶ Previous slides presented an **algorithm** for E-R  $\rightarrow$  relational mapping
  - ▶ The algorithm is definitely correct and straightforward to implement ...
  - ▶ But: it can generate many **redundant schema**
- ▶ What follows are some guidelines for *combining* or *removing* some of the redundant schema

## Weak Entity Sets



- ▶ The schema corresponding to a relationship set that links a **weak entity set** to its **identifying strong entity set** is redundant
  - ▶ Relationship *sec\_course* has no attributes of its own
  - ▶ The *section* schema already contains the primary-key attributes of the strong entity set (*course*) (see earlier discussion in today's lecture)
    - ▶ These (**plus the attributes of the *section* entity-set**) are precisely the attributes that would appear in the *sec\_course* schema
  - ▶ Every tuple in *sec\_course* has a matching tuple in *section* (and vice versa)
- ▶ Relationship schema is **redundant**: eliminate it!

## Take Advantage of Total Participation


- ▶ Reminder: if every entity  $e_1$  in  $E_1$  **must participate** in a relationship set with  $E_2$ 
  - ▶ We say that  $E_1$  has a **total participation** in that *many-to-one* relationship
- ▶ Proposal: instead of **three schemas**:  $E_1, E_2, R \dots$
- ▶ ...Let's combine  $E_1$  and  $R$  into a **a single schema**, leaving us with **two schema**
- ▶ Intuition: because of the **total participation** property, no need to keep  $E_1$  information separate from the relationship information
  - ▶ Note: the primary key of  $R$  is **already** the primary key of  $E_1$ 
    - ▶ Because  $E_1$  is on the “many” side of the relationship
  - ▶ So: this proposal doesn't add new attributes to primary key of  $E_1$
  - ▶ We only need to add relation-specific attributes (if any) to “enhanced”  $E_1$

## Total Participation $\Rightarrow$ Combining Schema

- ▶ Conceptually: merge  $R$  into  $E_1$
- ▶ Schema attributes of enhanced  $E_1$  are **union** of  $E_1$  and  $R$  attributes
  - ▶ Which is the union of the **primary key attributes** of  $E_1$  and  $E_2$  and **descriptive attributes** of  $R$  (if any)
- ▶ Primary key of enhanced  $E_1$  is primary key of original  $E_1$
- ▶ Only need one foreign-key constraint, to  $E_2$
- ▶ Result: we've eliminated the  $R$  schema 😊

- ▶ Note: we've previously introduced this “merge” approach (for **total participation**) in the context of 1 : 1 relationships!
- ▶ In that scenario, we can **arbitrarily select either** of the participating entity sets as the one that “assimilates the relationship set”
- ▶ The “merge” will not affect the primary key of the  $E_i$  into which the relationship is merged

## Combining Schema: Example




- ▶ This E-R diagram states that **every instructor must be associated with some department**
- ▶ Idea: instead of creating a schema for relationship set *inst\_dept* ...
  - ▶ Add the *dept\_name* attribute to the “enhanced” *instructor* schema
  - ▶ Because that’s the **primary key** of the *other entity set* (*department*) (the “1-side” of the 1 : N relationship)
- ▶ *instructor* attributes are now *ID*, *name*, *dept\_name*, *salary*
- ▶ *instructor* schema now has foreign-key constraint: attribute *dept\_name* references the *dept* schema

## Returning to Weak Entity Sets

- ▶ I hope that you now see that the approach for eliminating the **weak entity set relationship** is just a special case of the “total participation” approach that we just discussed
- ▶ In the case of a “weak entity set” we have ...
  - ▶ A many-to-one relationship
  - ▶ The weak entity-set has **total participation** in the relationship
  - ▶ There are no descriptive attributes in the relationship
  - ▶ The weak entity-set’s attributes already include the **primary key** attributes of the strong entity set

## Must We Have The Total Participation Property?

- ▶ Even if participation is **partial** on the “many” side, can still **combine schemas**!
- ▶ Trick:  $E_1$  (the “many side”) must store NULL values for all  $E_2$  primary key attributes whenever an entity in  $E_1$  has no mapping to an  $E_2$  entity



- ▶ Manager to employee mapping is one-to-many
- ▶ Relation schemas were
  - ▶ `employee(employee_id, name)`
  - ▶ `works_for(employee_id, manager_id)`
- ▶ Can combine into
  - ▶ `employee(employee_id, name, manager_id)`
  - ▶ But: Must store NULL for employees with no manager (the CEO))

# Merge Approach: Advantages & Disadvantages

- ▶ **Advantages** of combining schema:
  - ▶ Eliminate a foreign-key constraint (improve performance)
  - ▶ Eliminate an extra JOIN operation in certain queries
- ▶ **Disadvantages** of combining schema:
  - ▶ May require use of NULL values (partial participation case)
  - ▶ Makes it harder to change mapping cardinality constraints *in the future*



# Summary of E-R-Relational Mapping Constructs

## ER MODEL

Entity type

1:1 or 1:N relationship type

M:N relationship type

$n$ -ary relationship type

Simple attribute

Composite attribute

Multivalued attribute

Value set

Key attribute

## RELATIONAL MODEL

*Entity* relation

Foreign key (or *relationship* relation)

*Relationship* relation and *two* foreign keys

*Relationship* relation and  $n$  foreign keys

Attribute

Set of simple component attributes

Relation and foreign key


Domain

Primary (or secondary) key

Table copied from “Fundamentals of Database Systems”  
(Elmasri & Navathe)

# Review Of Common Relationship Set E-R $\Rightarrow$ RDB Mappings

# Many-to-many



## Product

<u>ID</u>	name
-----------	------

## Purchase

<u>prod_ID</u>	<u>person_ID</u>
----------------	------------------

## Person


<u>ID</u>	name
-----------	------

```
CREATE TABLE Product (  
  ...  
  PRIMARY KEY (ID)  
);
```

```
CREATE TABLE Purchase (  
  ...  
  PRIMARY KEY (prod_ID, person_ID),  
  FOREIGN KEY (prod_ID) REFERENCES Product (ID),  
  FOREIGN KEY (person_ID) REFERENCES Person (ID)  
);
```

```
CREATE TABLE Person (  
  ...  
  PRIMARY KEY (ID)  
);
```

## Many-to-many – a question



### Purchase

prod\_ID person\_ID


Primary key fields  
also NOT NULL.

Doesn't that force  
min. cardinality = 1?

```
CREATE TABLE Purchase (  
  ...  
  PRIMARY KEY (prod_ID, person_ID),  
  FOREIGN KEY (prod_ID) REFERENCES Product (ID),  
  FOREIGN KEY (person_ID) REFERENCES Person (ID)  
);
```

No. E.g., a **Product**  
never purchased  
simply never appears  
in **Purchase** table.

## One-to-many / many-to-one



Publishes table optimized away.

### Publisher

<u>ID</u>	name	address
-----------	------	---------


### Book

<u>ISBN</u>	title	num_copies	is_self_pub	pub_ID	pub_date
-------------	-------	------------	-------------	--------	----------

```
CREATE TABLE Publisher (  
    ...  
    PRIMARY KEY (ID)  
);
```

```
CREATE TABLE Book (  
    ...  
    PRIMARY KEY (ISBN),  
    FOREIGN KEY (pub_ID) REFERENCES Publisher (ID)  
);
```

## One-to-many / many-to-one – a confusion



### Publisher

<u>ID</u>	name	address
-----------	------	---------


### Book

<u>ISBN</u>	title	num_copies	is_self_pub	pub_ID	pub_date
-------------	-------	------------	-------------	--------	----------

Way to remember where the additional field is placed:

- **Book** has one publisher, so it can be a field.
- **Publisher** has many books, so it can't be a field.

# One-to-many with minimum cardinality



## Company

<u>ID</u>	name	address
-----------	------	---------

```
CREATE TABLE Company (  
  ...  
  PRIMARY KEY (ID)  
);
```

## Product

<u>ID</u>	name	co_ID
-----------	------	-------


```
CREATE TABLE Product (  
  ...  
  co_ID ... NOT NULL,  
  PRIMARY KEY (ID),  
  FOREIGN KEY (co_ID) REFERENCES Company (ID)  
);
```

Companies with no products don't appear.

Product w/o Company isn't allowed.

FK doesn't imply NOT NULL.

## Parallel one-to-many/one relationships



Same idea as before,  
for each relationship.

### Person

<u>ID</u>	name	address
-----------	------	---------

```
CREATE TABLE Person (  
  ...  
  PRIMARY KEY (ID)  
);
```


### Vehicle

<u>ID</u>	model	owner_ID	mechanic_ID
-----------	-------	----------	-------------

```
CREATE TABLE Vehicle (  
  ...  
  PRIMARY KEY (ID),  
  FOREIGN KEY (owner_ID) REFERENCES Person (ID),  
  FOREIGN KEY (mechanic_ID) REFERENCES Person (ID)  
);
```



# Weak entity



## University

name

```
CREATE TABLE University (  
  ...  
  PRIMARY KEY (name)  
);
```

## Team

univ\_name

ID


sport

```
CREATE TABLE Team (  
  ...  
  PRIMARY KEY (univ_name, ID),  
  FOREIGN KEY (univ_name) REFERENCES University (name)  
);
```

Primary key = keys  
of both entities.  
Must have value.

Foreign key = key  
of related table.  
Must have value.

## One-to-one with semantically same key




Usually indicates a poor design. Combine into one entity set.

### Player

<u>ID</u>	name	address	points
-----------	------	---------	--------

## One-to-one with different keys



Sometimes indicates a poor design.

### Faculty

<u>emp_ID</u>	name	office	dept_id
---------------	------	--------	---------

```
CREATE TABLE Faculty (  
  ...  
  PRIMARY KEY (emp_ID),  
  FOREIGN KEY (dept_ID) REFERENCES Dept (dept_ID)  
);
```

### Dept


<u>dept_ID</u>	name	chair_id
----------------	------	----------

```
CREATE TABLE Department (  
  ...  
  chair_id ... NOT NULL,  
  PRIMARY KEY (dept_ID),  
  FOREIGN KEY (chair_ID) REFERENCES Faculty (emp_ID)  
);
```

Same optimization  
as for one-to-many.

Generalizing to n-ary  
relationships


## N-ary with multiple “many” sides



Junction table:

- Primary key = combination of PKs of all “many”-side tables
- Foreign keys to all related entity sets' PKs


## N-ary with only one “many” side



Implemented same as multiple binary one-to-many relationships:

- No junction table.
- “Many”-side table has FKs to other tables.

N-ary with no “many” side



Again, depends on whether tables should be combined or not.

## Wrapping This Up

*A relational database is “basically” a set of tables with values drawn from specified domains, and represented using a relational schema*

I hope that you now appreciate that this technically correct definition basically misses the point

The following captures more of the essence of our discussion

- ▶ A relational database consists of
  - ▶ A set of tables with identifiers (primary keys)
  - ▶ A set of many-to-one binary relationships between tables, each of which is “induced” by foreign-key constraints
  - ▶ Each binary relationship represents a function (partial or total) from  $table_i \rightarrow table_j$



# Today's Lecture: Wrapping it Up

From E-R to Relational

Schema Redundancy

Review Of Common Relationship Set E-R  $\Rightarrow$  RDB Mappings

## Readings

- ▶ Today's lecture covered **Chapter 6.6 - 6.7** in the textbook
- ▶ Note: we didn't discuss "Extended E-R Features" (Chapter **6.8** in the textbook), nor did we discuss "Alternative Notations" (Chapter **6.10**)

- ▶ You are not responsible for this material, but as we leave this discussion, keep in mind that we definitely only "scratched the surface"
- ▶ We didn't even get into the topic of E-R design tools ☹
- ▶ At this point, you know enough to get by, and you know enough to be dangerous ☺

- ▶ Note: we've focused on **binary relationships**, and ignored the much more complicated issue of e.g., ternary relationships
  - ▶ One reason: not enough time ☹
  - ▶ Another reason: ternary (and higher) relationships can be converted into a set of binary relationships