

# Java Concurrency In a Nutshell

COM 3563: Database Implementation

Avraham Leff

Yeshiva University

*avraham.leff@yu.edu*

COM3563: Fall 2020

# Today's Lecture: Overview

1. Why A Lecture On “Java Concurrency”?
2. Transitioning To Multi-Threaded Code
3. Race Conditions: The Problem(s)
4. Synchronization: The Solution
5. Concurrency Programming Guidelines
6. Creating & Managing Multiple Clients
7. Some Closing Thoughts

- ▶ Note: we will likely not finish this material **in class**.
- ▶ Because this material is not “core” *Database Implementation* material, am not allocating more than one lecture
- ▶ For your (project implementation etc) sake: review (and think about) all the material

# Why A Lecture On “Java Concurrency”?

## How Is “Concurrency” Relevant To “Database Implementation”?

- ▶ As you probably know, this course is not *Parallel Programming* ☺
- ▶ Q: so, why a lecture on **Java concurrency**?
- ▶ A: because databases, like all “service providing” software, must perform well
  - ▶ Or: they won’t have any customers ☺
- ▶ Throughput: a fundamental performance metric
  - ▶ “How many client requests serviced per given time unit?”
- ▶ Implication: your DBMS must support **concurrent, multi-client access** to its data
- ▶ Your implementation must therefore address a fundamental tension
  1. Must allow concurrent access (“the more clients, the better”)
  2. Must ensure that concurrent access doesn’t corrupt data-structures (“the fewer clients, the easier to implement”)

## Why Java Concurrency?

- ▶ Having established that understanding concurrency issues & solutions is relevant to our course ...
- ▶ Why: Java concurrency?
  - ▶ After all: many (most?) enterprise databases are written in C or C++
  - ▶ (See [this discussion, for example](#))
- ▶ More importantly: the Java concurrency model is definitely different from C/C++
  - ▶ Java concurrency libraries are irrelevant to other languages ☹
- ▶ Frankly: was tempted to ignore the whole issue (as the textbook does ☺)
- ▶ But: you'll be implementing the course project in Java
  - ▶ To succeed, you must be competent in and comfortable with Java concurrency
  - ▶ Otherwise: your DBMS & my test cases will not get along ☹

You will have to teach yourself Java concurrency skills as you complete the course project

- ▶ This lecture will not teach you all or even much of what you need to know about Java concurrency
  - ▶ Consider taking *Parallel Programming* ☺
- ▶ Today’s lecture is a guided tour, tailored to helping you succeed in the course project
  - ▶ Make you aware of the pitfalls and some solution techniques
  - ▶ Steer you away from useless or incorrect information on the Internet

## Broader Importance Of This Material

- ▶ Any server-side software must allow for (and correctly manage) concurrency
- ▶ Keep in mind
  - ▶ Java is a **heavy-duty**, industrial-strength language
  - ▶ Its approach to concurrency is one of its **strengths**
  - ▶ Being able to do concurrent programming is a **key differentiator** between amateur and professional programmers
- ▶ I want you to incorporate these skills into your software portfolio
- ▶ Final point:
  - ▶ Yes: the solutions we'll discuss today are **specific to Java**
  - ▶ But: the issues in (shared memory) concurrent programming that we're discussing are **language independent**

## Historical Perspective (I)

- ▶ Primary goal of using **concurrency** is maximizing use of CPU (s)
- ▶ But: the strategy for maximizing CPU usage **has changed over time**
- ▶ **Single-core** era: use non-blocking I/O as an alternative to “blocking” code
  - ▶ Also: prioritized background tasks
- ▶ **Multi-core** era: Coarse-grained, task-based concurrency
  - ▶ Largely about **throughput**: pushing more requests through a server
  - ▶ That’s what we’ll focus on for this course
- ▶ **Many-core** era: Fine-grained data parallelism
  - ▶ Largely about latency: **“use more cores to get the answer faster”**

Credits to Brian Goetz, Java Language Architect for these insights



## Hardware Trends Drive Software Trends

- ▶ Hardware (previous slide) shapes the languages, library, and frameworks we write
- ▶ **Java 1**: supported threads, locks, condition queues
- ▶ **Java 5**: added thread pools, blocking queues, concurrent collections
- ▶ **Java 7**: added the fork-join library
- ▶ **Java 8**: added parallel streams

- ▶ You therefore want to focus on the **Java 5** additions
- ▶ `java.util.concurrent` and sub-packages (`atomic` and `locks`)

# Transitioning To Multi-Threaded Code

# Multi-Threading: Major Paradigm Shift (I)



- ▶ Since COM 1300, you've been “code-centric”
  1. You write code
  2. Computer executes that code “one line at a time”
- ▶ Your mind-set (for coding & debugging): *“I only have to worry about a single line of code at any given moment”*
  - ▶ Because the computer will execute **exactly one line of code at any given moment**

That mind-set is (mostly) **incorrect** ☺

# Multi-Threading: Major Paradigm Shift (II)

- ▶ What was really going on:
  1. You write code
  2. You create a **thread**: an independent execution “engine”
  3. That **thread** executes your code “one line at a time”
- ▶ Key point:
  - ▶ You can create multiple threads, and then direct each thread to **independently** execute that code
- ▶ Until now:
  - ▶ By default: the `java` command creates **only one thread to execute your program**
  - ▶ That thread is labeled as `main`
  - ▶ Typically, one thread suffices to get your work done, but not any more 😊

## Motivation: Performance! (I)

- ▶ I've already motivated this shift in your programming paradigm: **performance, performance, performance**
- ▶ Contrast to (for example) a *Data Structures* project
  - ▶ Your focus: write correct code that meets requirements
  - ▶ Elegant code is a bonus
- ▶ *Introduction to Algorithms* added the idea of “pick the right algorithm”
  - ▶ To break the “Big-O” barrier
  - ▶ But fundamentally: only one thread of computation is executing your code
- ▶ Now: your DBMS code is a “service”, available to **multiple, independent clients**
  - ▶ In other words: each of these “independent clients” is (in OS terms) a **thread** (or a process)
  - ▶ Because **OATsdb** is an **embedded database**, each of these “independent clients” is a **Java thread** (not a process)

## Motivation: Performance! (II)

- ▶ Theoretically: your DBMS could insist “I’ll service only one client at a time”
  - ▶ This would definitely make it much easier to write “correct & elegant code” 😊
  - ▶ But: your **throughput metric** will be very disappointing

- ▶ Implication

- ▶ Your development test-bed will have to create and manage **multiple clients** (threads)
  - ▶ Your DBMS implementation will have to ensure that these multiple threads don’t corrupt your data-structures

- ▶ We say: “*your DBMS must support **concurrency***”

## Key “Take-Aways” At This Point

- ▶ Stop thinking of your code as a “live” entity that can do **only one thing at a given moment**
  - ▶ Example: a *Stack* which can do either `pop` or `push` **but not both** at the same time
- ▶ Instead: your code is a “static” entity that is “infused with life” by **being executed by a Thread**
- ▶ If multiple threads execute your code, any number of your code points can be executed at the same time
  - ▶ Example: `Stack.pop` **and** `Stack.push`
- ▶ Murphy’s Law says: the worst possible combination of concurrent executions will occur 😊

## Remainder Of Lecture

1. Explain why multi-threaded programs require a paradigm shift on your part
  - ▶ In other words: the “concurrency problem”
2. Present solutions to the “concurrency problem”
3. Turn from the “server-side” DBMS issues to “client-side” test-bed issues



# Race Conditions: The Problem(s)

## Bad Interleavings: Method Boundaries May Not Preserve Invariants (I)

Since COM 1300 you've become accustomed to rely on the notion of **invariants** to establish program correctness

*In computer programming, specifically object-oriented programming, a class invariant (or type invariant) is an invariant used for constraining objects of a class. Methods of the class should preserve the invariant. The class invariant constrains the state stored in the object.*

*Class invariants are established during construction and constantly maintained between calls to public methods. **Code within functions may break invariants as long as the invariants are restored before a public function ends.***

**Source: Wikipedia**

## Bad Interleavings: Method Boundaries May Not Preserve Invariants (II)

- ▶ Once we enter the world of multi-threaded programming, we can no longer rely on invariants to establish program correctness
- ▶ In reality, methods are not executed atomically
- ▶  $Thread_1$ , executing  $method_1$ , may have its execution interleaved with  $Thread_2$ , executing  $method_2$
- ▶ Implication
  - ▶  $thread_1$  may see “intermediate” state affected by  $method_2$  execution
  - ▶  $thread_2$  may see “intermediate” state affected by  $method_1$  execution

## Bad Interleavings: Shared State Transitions Are Now Exposed

- ▶ The whole point of “object-oriented” programming is to **encapsulate shared state**
- ▶ Concurrent thread executions imply that (unless we're careful) we've broken encapsulation ☹
- ▶ Typical scenarios of different methods accessing **shared state**
  - ▶ One thread deposits money in a bank-account, other thread does a withdrawal
  - ▶ “Producer” thread enqueues a task in a queue, “consumer” thread dequeues a task
  - ▶ One thread does a *Hashtable.get*, other thread does a *Hashtable.put*
- ▶ Method abstraction conceals the fact that e.g., a single *Hashtable.get* is “really” multiple lower-level machine instructions
  - ▶ And thus can be interleaved with the execution of *Hashtable.put*
  - ▶ With disastrous results ☹

# Your Intuition Needs an Upgrade

This code is absolutely correct in a single-threaded environment

```
1  class BankAccount {
2      private int balance = 0;
3      int getBalance() { return balance; }
4      void setBalance(int x) { balance = x; }
5      void withdraw(int amount) {
6          int b = getBalance();
7          if (amount > b) {
8              throw new WithdrawTooLargeException();
9          }
10
11         setBalance(b - amount);
12     }
13
14     // other operations like deposit, etc.
15 } // class BankAccount
```

## Multi-Threaded Environment: Interleaving

Scenario ( $x$  and  $y$  are *BankAccount* instances):

- ▶ Thread  $T_1$  calls `x.withdraw(100)`
  - ▶ Thread  $T_2$  calls `y.withdraw(100)`
- 
- ▶ If  $T_1$  and  $T_2$  executions overlap (to **any** extent) , we say the executions **interleave**
  - ▶ Because of OS “time-slicing”, can happen even with a **single processor**
  - ▶ No problem if the threads are accessing **different** account instances
  - ▶ But: if the threads access the same account, we have **absolutely no idea what's going to happen**

## Interleaving: One Scenario

### Thread 1

```
int b = getBalance();  
  
if (amount > b)  
    throw new ...;  
setBalance(b - amount);
```

### Thread 2

```
int b = getBalance();  
if (amount > b)  
    throw new ...;  
setBalance(b - amount);
```

- ▶ Both withdrawals are successful, customer gets to “double-dip”
- ▶ This is a “lost-withdrawal” scenario, bank is very unhappy 😞

## Rearranging Code Doesn't Work

```
1 void withdraw(int amount) {  
2     if (amount > getBalance()) {  
3         throw new WithdrawTooLargeException();  
4     }  
5     setBalance(getBalance() - amount);  
6 }
```

- ▶ Rearranging or repeating code **does not solve** interleaving problems
- ▶ Here: only **narrows** the exposure, but other thread can still change the balance between second call to *getBalance* and reducing by *amount*
  - ▶ May not even do that 😊...since compiler may optimize back to original version
- ▶ May even change to “negative balance” scenario
  - ▶ Second call to *getBalance* occurs **after** other thread has already reduced balance
- ▶ Again: you simply can't assume that this won't happen



## “Race” Conditions (I)

A **race condition** occurs when the computation result depends on scheduling: specifically, how hardware & software interact to **interleave the concurrent execution of threads**

- ▶ I’m only introducing the term “race condition” because it’s so prevalent
- ▶ Unfortunately, it’s an overloaded term
  - ▶ It refers to **“bad” interleavings** (the problem we’ve been discussing so far)
  - ▶ And it refers to **data races**: a simultaneous read/write or write/write (by two threads) of the **same memory location**
- ▶ The common denominator of these two types of race conditions is that they’re **concurrency bugs**
  - ▶ If only one thread, then no problem manifests

## “Race” Conditions (II)

The good news: we’ll use the same solution approach to both “bad interleavings” and “data races”

- ▶ (I’m therefore going to postpone discussion of “data race” problem until we’ve discussed this solution)
- ▶ Conceptual difference between “data race” and “bad interleaving” is that a “data race” condition is always an error!
- ▶ In contrast: only you (the programmer) can determine whether an interleaving is “bad”
- ▶ Depends on how you’ve defined your **invariants** (the “code specification”)
  - ▶ “Bad interleaving” exists *iff* your code exposes an **intermediate state that violates an invariant**

## Segue: Solution For Race Conditions Is Mutual Exclusion

- ▶ Think of it this way: our concurrency problems are caused by allowing **independent threads to run wild** in our code
- ▶ Solution idea: impose restrictions on thread executions
  - ▶ Specifically: annotate our code to state *“only one thread can execute this block of code at a time”*
- ▶ Example: **allow at most one thread to withdraw from account at a time**
  - ▶ Note: this requires that we **also exclude** other concurrent operations such as *deposit*
- ▶ This technique is called **mutual exclusion** or the creation of **critical sections**



# Support For Mutual Exclusion

- ▶ Unfortunately: mutual exclusion **cannot be enforced by the compiler** ☹
  - ▶ The semantics are **at a higher-level** than the compiler understand
  - ▶ How can the compiler possibly understand what interleavings are OK and which ones are dangerous?
- ▶ We do need, and can get, support from Java language primitives
  - ▶ The **language must be able to enforce** mutual exclusion

## One Approach: Locks or Mutex Variables

- ▶ Some languages and **libraries** use the concept of lock or **mutex variable** to provide mutual exclusion
- ▶ Conceptually: two methods
  - ▶ acquire
  - ▶ release
- ▶ A critical section (such as accessing an account's balance variable) is ...
- ▶ ...Preceded by a call to acquire (*thread<sub>2</sub>* will be forced to **block** if *thread<sub>1</sub>* **currently holds the lock**)
- ▶ ...Followed by a call to release (*thread<sub>1</sub>* will release the lock after it modifies balance, allowing *thread<sub>2</sub>* to acquire the lock and proceed)
- ▶ Locks require **special hardware and os support** to implement
  - ▶ We're requiring that "*check if lock is held and if not held acquire the lock*" be an **atomic operation**

## Lock Variables: Getting It Wrong

- ▶ Wrong: if you use different locks for `withdraw` and `deposit`
  - ▶ Mutual exclusion works only when the shared resource being protected (`balance` instance variable) is protected by the same lock
- ▶ Bad Performance: too coarse a lock granularity
  - ▶ Example: don't use one lock for all bank accounts 😊
- ▶ Wrong: if `thread1` “forgets” to release a lock, then it will block other threads forever
- ▶ Example:
  1. Thread acquires lock
  2. Thread invokes `withdraw`
  3. Withdrawal attempt triggers `WithdrawTooLargeException`
  4. The lock will never be released
  5. All threads are prevented from using this code forever

## Locks: Reentrantcy Problem

- ▶ Assume you've used locks correctly: `withdraw` and `deposit` use the **same lock**
- ▶ Now you implement `getBalance` and `setBalance`
- ▶ Two scenarios
  1. The new code uses a **different lock** than the “`withdraw`” lock
    - ▶ Incorrect: you've enabled a race condition between `setBalance` and `withdraw`
  2. The new code uses the **same lock** as the “`withdraw lock`”
    - 2.1 `withdraw` acquires the lock
    - 2.2 `withdraw` (internally) invokes `getBalance`
    - 2.3 Thread **blocks forever** as it tries to acquire the `getBalance` which **it already has** 😊
  3. One solution: allow locks to be **reentrant**



## Reentrant Locks: Semantics

- ▶ A **reentrant lock** (sometimes called a *recursive lock*)
  - ▶ “Remembers” which thread has acquired it
  - ▶ Maintains a `count` variable
- ▶ When lock goes from “not held” to “acquired” state ...
  - ▶ `count` is set to 1
- ▶ If acquiring thread (re)invokes “acquire” on this lock ...
  - ▶ Thread **does not block**
  - ▶ Instead: `count` is incremented
- ▶ When thread invokes “release” on this lock ...
  - ▶ `count` is decremented
  - ▶ If `count` now reaches 0, lock enters “not held” state

## Segue: Java Object Monitors

- ▶ I've included this discussion of the **lock approach** because
  - ▶ Many languages and libraries use it
  - ▶ Because this is the approach used in the JDK 5 `java.util.concurrent.locks.ReentrantLock` class(es) ☺
  - ▶ Because it will help you understand the **Java language solution: object monitors**
- ▶ My opinion only: the Java object monitor approach is superior for “vanilla” mutual exclusion
  - ▶ You can't beat built-in language support
  - ▶ Specifically “**automatic acquire and release**”
- ▶ Use the JDK 5 facilities for more complicated situations
  - ▶ And read the Javadoc carefully!

# Java Object Monitors: Introduction

- ▶ In Java, every object instance and every class is associated with a monitor
  - ▶ A monitor is implicitly associated with a lock
  - ▶ Java manages the magic transparently
- ▶ When a thread acquires the monitor lock, all other threads are blocked from acquiring that lock
- ▶ How does a thread acquire the monitor lock?
  - ▶ Not by accessing the object's code!
  - ▶ By (literally) synchronizing on the object instance

```
1 synchronized (o) {  
2     // code goes here  
3 }
```

- ▶ We say that the thread has synchronized on Object o
- ▶ JVM will block this thread if some other thread currently has “Object o lock”

# Java Synchronization

```
1    synchronized (o) {  
2        // code goes here  
3    }
```

- ▶ Thread has the lock if it makes it past beginning of synchronized block
- ▶ JVM will automatically release that lock when thread executions moves past the **closing curly brace** of the synchronized block
  - ▶ Even if it exits via an Exception 😊
- ▶ Note (I can't stress this too much): complete decoupling of the **thread** ...
  - ▶ A “live thing”, something which **executes code**
- ▶ From the **lock** ...
  - ▶ Which is associated with an **object (or class) instance**
  - ▶ Static code and associated mutable state

# “Synchronization” Is Just a Lock

- ▶ I introduced Java **synchronized** keyword after discussing **reentrant lock approach** because
  - ▶ That’s essentially what a Java monitor is: a reentrant lock associated with an object or class instance
  - ▶ For which the language and JVM support provide automatic “acquire” and “release” semantics
- ▶ But conceptually the same
- ▶ With all the perils thereto 😊
  - ▶ **You are responsible** for ensuring that the lock is used consistently to protect the relevant critical section **throughout your code**

## Synchronize On Private Object?

You can synchronize on explicit lock variable

```
1 class BankAccount {
2     private int balance = 0;
3     private Object lk = new Object();
4     // etc
5
6     void withdraw(int amount) {
7         synchronized(lk) {
8             // modify balance
9         }
10    }
11 }
```

But ...

- ▶ Makes it harder for other code to synchronize on this BankAccount's operations
- ▶ And ...why bother? Just synchronize on **this**
  - ▶ The BankAccount instance itself

```
1 class BankAccount {
2     private int balance = 0;
3     int getBalance() { synchronized (this) { return balance; }
4     }
5 }
```

# Can Synchronize On Method Granularity

```
1    class BankAccount {  
2        private int balance = 0;  
3        synchronized int getBalance() { return balance; }  
    }
```

- ▶ Syntactic sugar ...
- ▶ Putting `synchronized` before a method declaration is equivalent to surrounding the method body with a `synchronized(this)`
- ▶ Looks nicer 😊
- ▶ May cost you if the method consists of multiple, “long-duration” critical sections involving `different` shared resources
  - ▶ Prevents other threads from accessing all of those resource

## “Data Races”

- ▶ In addition to solving the “bad interleaving” problem, mutual exclusion solves the “data race” problem
- ▶ Key point: sensitize yourself to the existence of such problems

*What is a Data Race?*

*The Thread Analyzer detects data-races that occur during the execution of a multi-threaded process. A data race occurs when:*

- 1. two or more threads in a single process access the same memory location concurrently, and*
- 2. at least one of the accesses is for writing, and*
- 3. the threads are not using any exclusive locks to control their accesses to that memory.*

*When these three conditions hold, the order of accesses is non-deterministic, and the computation may give different results from run to run depending on that order. Some data-races may be benign (for example, when the memory access is used for a busy-wait), but many data-races are bugs in the program.*



## Data Race: This Code Is Incorrect

```
1  class Stack<E> {  
2      private E[] array = (E[]) new Object[SIZE];  
3      int index = -1;  
4      boolean isEmpty() { // unsynchronized: wrong?  
5          return index == -1;  
6      }  
7  
8      synchronized void push(E val) {  
9          array[++index] = val;  
10     }  
11  
12     synchronized E pop() {  
13         return array[index--];  
14     }  
15  
16     E peek() { // unsynchronized: wrong!  
17         return array[index];  
18     }  
19 }
```

- ▶ Perhaps *isEmpty* and *peek* can “get away” with dropping **synchronized**?
  - ▶ After all: looks like the state changes written or read by these methods are atomic

## Data Race Problems: You Must Synchronize

- ▶ Even “tiny steps” may require multiple steps in the compiled code implementation
  - ▶ `array[++index] = val` (probably) takes at least two steps
- ▶ Useful to think of a “data race” as a “simultaneous access error”
- ▶ Compared to “bad interleavings”, a “data race” is a “lower-level” error
  - ▶ Absolutely, fundamentally, a problem
  - ▶ You simply cannot reason about the behavior of code that has data races
  - ▶ From the C++ reference: “If a data race occurs, the behavior of the program is undefined”
- ▶ Not going into **how data races actually cause problems**
  - ▶ That’s a discussion for a compiler and/or computer architecture course 😊
- ▶ See the `java.util.concurrent.atomic` classes if you want “atomic” operations on individual variables

## A Word About Deadlocks


- ▶ Students' reaction to discussion so far is often *“no problem, I'll wrap all my methods with `synchronized`”* 😊
- ▶ One issue: may hurt performance because other threads can't make progress while this thread is executing a given method
  - ▶ They can execute neither this method nor any other method of this object
  - ▶ We'll elaborate on this issue later
- ▶ More fundamental problem: **deadlock**
  1.  $Thread_1$  acquires the lock but blocks before releasing (e.g., “data structure is full”)
  2.  $Thread_2$ , waiting to empty the data structure, can't acquire the lock
  3. Both threads block forever
- ▶ Classic example: a **“producer/consumer”** class
- ▶ Lesson: don't blindly “synchronize” everything



# Introduction

- ▶ Previous material is about making you aware of the **multi-threading** paradigm shift
- ▶ Why perfectly correct sequential code can break as soon as it's deployed in a multi-threaded environment
- ▶ The need for using mutual exclusion techniques to keep concurrent threads from stepping on one another
- ▶ Now: some guidelines as you begin writing your own concurrent code

## Categorize Memory Into Three Buckets



Every memory location in your program should have (at least) one of the following properties

- ▶ Be **thread-local**: only one thread ever accesses it
- ▶ Be **immutable**: (after being initialized), memory is only read, **never written**
- ▶ Be **synchronized**: locks are used to ensure there are no race conditions

## Thread-Local Memory

- ▶ Simplest way to avoid data races is for only one thread to have access to memory 😊
- ▶ So: whenever possible, **do not share resources among threads**
- ▶ Definitely encourage you to read the `java.lang.ThreadLocal` Javadoc!
- ▶ If multiple threads need to access a resource, see if you can give each a **copy of that resource**
  - ▶ Example: `java.util.Random`
  - ▶ Only consider allowing threads to **share a mutable resource** if no other choice
- ▶ Note: this is why you don't need to synchronize on a method's **local variables** ...
- ▶ The runtime's call-stack is (effectively) thread-local!

# Immutable Memory

- ▶ Whenever possible, do not **update objects**
  - ▶ Instead: **create new objects**
- ▶ This is one of the key tenets of **functional programming**
  - ▶ Always (even in **sequential environment**) a good idea to avoid side-effects
  - ▶ An invaluable idea in a **concurrent environment**!
- ▶ Key point: concurrent **read operations** can never cause data races
- ▶ Programming tip: we're biased to "reusing" (**mutating**) object instances
  - ▶ We justify that bias by talking about "efficiency" ☺
  - ▶ Try to correct for that bias



## The Rest of The Venn Diagram

- ▶ After you've minimized the amount of state **shared & mutable state**, the burden of providing mutual exclusion falls on you
- ▶ Some guidelines for these scenarios ...
- ▶ Guideline #0: No data races
  - ▶ Never allow two threads to read/write or write/write the same location at the same time
- ▶ (Remember: avoiding “data races” **will not necessarily fix** “bad interleavings”)
- ▶ Guideline #1: For each piece of state that requires mutual exclusion, associate a lock that must be acquired before thread can read or write that location
  - ▶ We say the lock **guards the location**


# Consistent Locking

- ▶ Perfectly fine to use the same lock to guard multiple pieces of state
  - ▶ In Java, the guard is often the object containing the location
  - ▶ `this` inside an object's methods
  - ▶ But you can guard a larger data-structure with its own lock to provide more granular mutual exclusion
- ▶ Remember to document the guard for each location
- ▶ You – the programmer – get to decide how to partition the control of your “shared-and-mutable” state lock

## Flexibility: Change The Locking Protocol Dynamically

- ▶ We already have three techniques to avoid problems
  1. Thread-local
  2. Immutable
  3. Consistent locking
- ▶ Your code can dynamically switch between one technique and another
  - ▶ As long as all threads agree on the protocol being used at any given time
- ▶ Example: data-structure initialization may use locking technique because multiple threads are used for performance reasons
  - ▶ Then kill all threads but “main”: now protected by “thread-local” technique
  - ▶ Or: once initialized, data-structure becomes immutable

## “Coarse-Grained” Versus “Fine-Grained” Locking



- ▶ Fewer locks, or “more objects per lock”
  - ▶ Example: One lock for entire data structure (e.g., array)
  - ▶ Example: One lock for all bank accounts



- ▶ More locks, or “fewer objects per lock”
  - ▶ Example: One lock per data element (e.g., array index)
  - ▶ Example: One lock per bank account

**Not a dichotomy:** this distinction is a continuum!

# “Coarse-Grained” Versus “Fine-Grained”: Tradeoffs

- ▶ **Coarse-grained advantages**
  - ▶ Simpler to implement
  - ▶ Faster/easier to implement operations that access multiple locations
    - ▶ Because all guarded by the same lock
- ▶ **Fine-grained advantages**
  - ▶ Increase possibilities for concurrent access
  - ▶ This can improve performance in situations where coarse-grained locking leads to unnecessary blocking
- ▶ **Guideline #2: Start with coarse-grained (simpler) and move to fine-grained (performance)**  
only if contention on the coarser locks becomes an issue

---

  - ▶ Use a profiler before trusting your intuition!
  - ▶ Can guarantee that (at least initially) you'll introduce lots of bugs 😊

## “Critical-Section” Granularity

- ▶ We’ve been discussing **lock granularity**: how many objects per lock
  - ▶ Orthogonal issue: **critical-section granularity**
    - ▶ How much “work” in the code that must be done while holding locks
  - ▶ Critical sections are too long?
    - ▶ Performance loss because other threads are blocked
  - ▶ Critical sections are too short?
    - ▶ Bugs! You’ve written code where other threads will be able to see **intermediate state**
  - ▶ Guideline #3: Make an effort to **not do expensive computation or I/O in critical sections**
    - ▶ And take even more care not to introduce race conditions
- ☺

# Understanding The “Critical-Section Granularity” Issue

- ▶ Tempting (and often a good idea) to simply wrap a chunk of code with a single lock
  - ▶ Example: code only contains operations like assignment statements and “cheap” method
    - ▶ Code runs so fast that it’s not worth splitting into multiple critical sections
- ▶ Assumption fails when critical section will contain “expensive” computation
- ▶ Scenario: you need to **replace** a Hashtable **value** by performing some computation on its **old value**
  - ▶ We want other threads to see **either** the old value **or** the new value
  - ▶ We also need the old value **in order to perform the expensive code** required to compute the new value

## This Critical-Section Is Too Large

- ▶ Assume that `lock` guards the Hashtable
- ▶ The code below locks the entire table during the expensive computation

```
1    synchronized(lock) {  
2        v1 = table.get(k);  
3        v2 = expensive(v1);  
4        table.put(k, v2);  
5    }
```



## This Critical-Section Is Too Short

```
1    synchronized(lock) {  
2        v1 = table.get(k);  
3    }  
4  
5    v2 = expensive(v1);  
6  
7    synchronized(lock) {  
8        table.put(k, v2);  
9    }
```

- ▶ The above code is buggy!
- ▶ Scenario: *thread*<sub>2</sub> updates to map  $k \rightarrow v_3$  while *thread*<sub>1</sub> was executing `expensive(v1)`
- ▶ If *thread*<sub>2</sub> does the update **after** *thread*<sub>1</sub>, then table will end up with  $k \rightarrow v_3$  (**not v2**)
- ▶ If *thread*<sub>2</sub> does the update **before** *thread*<sub>1</sub>, then table will end up with  $k \rightarrow \text{expensive}(v_3)$  (**not v2**)
- ▶ Desired semantics:  $k \rightarrow v_2$

## This Critical-Section Is Just Right

```
1  boolean loop_done = false;
2  while(!loop_done) {
3      synchronized(lock) {
4          v1 = table.get(k);
5      }
6
7      v2 = expensive(v1);
8
9      synchronized(lock) {
10         // If true, we already have correctly
11         // computed ‘v2’
12         //
13         // If false, we know that some other
14         // thread has updated the map: throw out
15         // pre-computed value and start over
16         if (table.get(k) == v1) {
17             loop_done = true;
18             table.put(k, v2);
19         }
20     } // while-loop
```

Note: this approach assumes that **it doesn't matter** if some arbitrary sequence of Hashtable operations occurred between the two critical section. All we care about is that we change mapping from  $k \rightarrow v1$  to  $k \rightarrow v2$

# Design For Atomicity

- ▶ Operations on your data-structures should be **atomic**
  - ▶ No other thread can see the “partial execution” of that operation
- ▶ Guideline #4: Think about atomicity **first** and locks **second**
  - ▶ “What operations need to be atomic?”
  - ▶ Make critical sections just long enough to preserve atomicity
  - ▶ **Then** design the locking protocol to implement the critical sections correctly

## Don't "Roll Your Own" (I)

- ▶ You need to understand today's material because you'll be writing **new code**
- ▶ But: make sure you're aware of the JDK 5 `jdk.util.concurrent` **concurrent collection** APIs
  - ▶ These are high-performance concurrent implementations of standard collection interfaces such as *List*, *Queue*, and *Map*
  - ▶ To provide high concurrency, these implementations manage their own synchronization internally
  - ▶ Implication: you can't exclude concurrent activity from a concurrent collection
  - ▶ Locking a concurrent collection only slows your code down, accomplishes nothing!
- ▶ Take the time to understand what these classes do and how they do it **before coding** your own data-structure
  - ▶ See if you can re-purpose the concurrent collection classes even if it doesn't immediately look like it addresses your problem

## Don't "Roll Your Own" (II)

```
1  // Code by Bloch, Item #81: "Simulate the
2  // behavior of String.intern:"
3  private static final ConcurrentMap<String, String> map =
4      new ConcurrentHashMap<>();
5
6  public static String intern(String s) {
7      // ConcurrentHashMap is optimized for
8      // retrieval operations, such as
9      // get. Therefore, it is worth invoking
10     // get initially and calling putIfAbsent
11     // only if get indicates that it is
12     // necessary:
13     String result = map.get(s);
14     if (result == null) {
15         result = map.putIfAbsent(s, s);
16         if (result == null)
17             result = s;
18     }
19     return result;
20 }
```

## Concurrent Collections: Bloch's Observations

- ▶ (FWIW, my own performance tests confirm Bloch's statements below)
- ▶ “On my machine, the `intern` method above is over six times faster than `String.intern`”
- ▶ “Concurrent collections make synchronized collections largely obsolete”
- ▶ “For example, use `ConcurrentHashMap` in preference to `Collections.synchronizedMap`”
- ▶ “Simply replacing synchronized maps with concurrent maps can dramatically increase the performance of concurrent applications”

## Blocking Operations

- ▶ In JDK 5, some of the collection interfaces were extended with blocking operations
  - ▶ That is: if a method can't be performed **at the time that a thread invoked it**...
  - ▶ The thread will (automatically) **block** until the internal state of the collection permits the method to continue
- ▶ Example: `BlockingQueue` extends `Queue` and adds several methods, including `take`
  - ▶ `take` removes and returns the head element from the queue, waiting if the queue is empty
- ▶ Blocking queues can be used to implement **work queues** (also known as “**producer-consumer**” queues)
  - ▶ One or more **producer** threads enqueue work items
  - ▶ One or more **consumer** threads dequeue and process items as they become available
  - ▶ Threads block as long as queue state isn't appropriate for desired semantics

# Synchronizers

- ▶ **Synchronizers** are objects that enable threads to **wait for one another**
  - ▶ Use synchronizers when you need to **coordinate** decoupled activities
- ▶ Typically, the *CountDownLatch* and *Semaphore* classes suffice
  - ▶ Sometimes you need the more powerful (and complicated) *CyclicBarrier* and *Exchanger* classes
  - ▶ Or even the *Phaser*



## CountDownLatch (I)

- ▶ Countdown latches are “single-use” barriers
  - ▶ They allow one or more threads to wait for one or more other threads to do something else
  - ▶ CountDownLatch constructor takes an `int`: specifies the number of times the `countDown` method must be invoked on the latch before all waiting threads are allowed to proceed
- ▶ Scenario (see next slide for code)
  - ▶ You want to “time” code execution for a **given “concurrency level”**
  - ▶ That implies that you only start the clock when all worker threads are “ready to go”
  - ▶ Implementation: one latch is initialized to the given concurrency level, each worker decrements the count when ready
  - ▶ That latch prevents timer from starting until all are ready
  - ▶ Another latch ensures that the timer measurement is only taken when all workers have completed

## CountDownLatch (II)

```
1 // Code from Goetz and Bloch
2 public class ConcurrentTimer {
3     private ConcurrentTimer() { } // Noninstantiable
4
5     public static long time(Executor executor, int concurrency,
6                             Runnable action) throws InterruptedException {
7         CountDownLatch ready = new CountDownLatch(concurrency);
8         CountDownLatch start = new CountDownLatch(1);
9         CountDownLatch done = new CountDownLatch(concurrency);
10
11         for (int i = 0; i < concurrency; i++) {
12             executor.execute(() -> {
13                 ready.countDown(); // Tell timer we're ready
14                 try {
15                     start.await(); // Wait till peers are ready
16                     action.run();
17                 } catch (InterruptedException e) {
18                     Thread.currentThread().interrupt();
19                 } finally {
20                     done.countDown(); // Tell timer we're done
21                 }
22             });
23         }
24
25         ready.await(); // Wait for all workers to be ready
26         long startNanos = System.nanoTime();
27         start.countDown(); // And they're off!
28         done.await(); // Wait for all workers to finish
29         return System.nanoTime() - startNanos;
30     }
31 }
```

# Creating & Managing Multiple Clients

# Introduction

- ▶ Old-style discussions of how to create multiple, independent, tasks focus on “*should I extend `Thread` or should I implement `Runnable`?*”
- ▶ That’s really “old-style”: almost always, **refrain from working directly with threads!**
- ▶ When you work directly with threads, a `Thread` serves as both a “**unit of work**” (or “**task**”) and the **mechanism for executing** the unit of work
- ▶ JDK 5 provides the *Executor* framework, which separates the **task** concept from the **execution mechanism**
- ▶ Two kinds of tasks: *Runnable* and its close cousin, *Callable*
  - ▶ A *Callable* is like *Runnable*, except that it **returns a value** and **can throw arbitrary exceptions**

## Executor Service

- ▶ The general mechanism for executing tasks is the JDK 5 (interface-based) **executor service**
- ▶ Think in terms of tasks (not Threads!), let an executor service execute those tasks for you
  - ▶ You'll gain the flexibility of scheduling different execution policies, concurrency levels, and other tunable knobs
  - ▶ Read the Javadoc!
- ▶ Some examples:
  - ▶ You can wait for any or all of a collection of tasks to complete (see `invokeAny` and `invokeAll`)
  - ▶ You can wait for the executor service to terminate (see `awaitTermination`)
  - ▶ You can retrieve the results of tasks one by one as they complete (using an *ExecutorCompletionService*)
  - ▶ You can schedule tasks to run at a particular time or to run periodically (using a *ScheduledThreadPoolExecutor*)

# Executor Example

```
1  // Code from Java Concurrency In Practice (Goetz)
2  public class PrimeGenerator implements Runnable {
3      private static ExecutorService exec = Executors.newCachedThreadPool();
4      private final List<BigInteger> primes = new ArrayList<BigInteger>();
5      private volatile boolean cancelled;
6
7      public void run() {
8          BigInteger p = BigInteger.ONE;
9          while (!cancelled) {
10             p = p.nextProbablePrime();
11             synchronized (this) {
12                 primes.add(p);
13             }
14         }
15     }
16     public void cancel() {
17         cancelled = true;
18     }
19     public synchronized List<BigInteger> get() {
20         return new ArrayList<BigInteger>(primes);
21     }
22     static List<BigInteger> aSecondOfPrimes() throws InterruptedException {
23         PrimeGenerator generator = new PrimeGenerator();
24         exec.execute(generator);
25         try {
26             SECONDS.sleep(1);
27         } finally {
28             generator.cancel();
29         }
30         return generator.get();
31     }
32 }
```

## Some Closing Thoughts

# “Concurrent” Programming

- ▶ Concurrency: the problem of **correctly** and **efficiently** managing concurrent access to shared resources from multiple clients
- ▶ Requires coordination between threads
  - ▶ Typically done by requiring other threads to **block** until “first thread is no longer accessing the shared resource”
  - ▶ Compare to the “fork/join” framework from *Introduction to Algorithms*
    - ▶ Concurrency coordination is very different from “fork/join” coordination
    - ▶ *join* semantics:  $thread_1$  blocks until  $thread_2$  has **completely finished its execution**
    - ▶ Concurrency coordination:  $thread_1$  blocks until  $thread_2$  has finished some “arbitrary” section of code execution
- ▶ Major challenge: correct concurrent programs are **non-deterministic**
  - ▶ Non-repeatability → much more complicated to test & debug than single-threaded environments ☹



## Parallelism Versus Concurrency (I)

- ▶ We've previously discussed (in *Introduction to Algorithms*) using **multiple threads** to achieve **parallelism**
  - ▶ Defined as “how to break up a computation into parts that can be executed in parallel”
  - ▶ Remember the **Fork/Join framework**?
- ▶ These parallel algorithms all had a very simple structure that avoided the **race conditions** discussed in today's lecture
- ▶ Each thread is assigned its “own section of memory”
  - ▶ Example: algorithm partitions array into sub-arrays, one sub-array per thread
- ▶ Yes: all threads have access to the same shared memory (“entire array”)
  - ▶ But algorithm ensures that threads agree as to what's “mine” and what's “theirs”
- ▶ On `fork`, algorithm completely assigns ownership of sub-array to “forkee”
- ▶ Main thread only **resumes ownership** when it does a `join` on the “forkee” thread

## Parallelism Versus Concurrency (II)

- ▶ Unfortunately for us, the approach we used for fork/join **parallelism** won't work when
  - ▶ Memory is accessed by threads in overlapping or unpredictable ways
  - ▶ Threads are doing **un-coordinated tasks** but still need concurrent access to the **same resources**
  - ▶ Such scenarios raise the **mutual-exclusion** issues discussed in today's lecture
- ▶ Thought experiment: can you “reduce” implementation of a data-structure (let alone a DBMS) to fork/join parallelism?
- ▶ In a nutshell:
  - ▶ **Parallelism** is about making algorithms run faster
  - ▶ **Concurrency** will not make “an algorithm” run faster

# Concurrency: Not Just About Throughput

- ▶ We've previously explained that concurrency is about achieving **good throughput**
  - ▶ Alternatively: enabling as many client threads to use a service (or data-structure) at the same time (while preserving "**correctness**")
- ▶ But note that concurrency also improves another performance metric: **latency**
- ▶ No time to pursue this idea, but here are some examples
  - ▶ Respond to GUI events in one thread while another thread performs an expensive computation
  - ▶ Web-application services another request while another thread does I/O for previous request
  - ▶ Failure in  $task_1$  does not bring down  $task_2$

# Today's Lecture: Wrapping it Up

Why A Lecture On “Java Concurrency”?

Transitioning To Multi-Threaded Code

Race Conditions: The Problem(s)

Synchronization: The Solution

Concurrency Programming Guidelines

Creating & Managing Multiple Clients

Some Closing Thoughts

- ▶ I've already strongly recommended [Java Concurrency in Practice](#) by *Brian Goetz et al* [Amazon](#)
- ▶ This lecture overlaps considerably with Chapters 2 & 3
- ▶ Am not assigning “homework”, just a course project ☺