

# More “Basic” SQL: Set Operations & NULL Values

COM 3563: Database Implementation

Avraham Leff

Yeshiva University

*avraham.leff@yu.edu*

COM3563: Fall 2020





- ▶ We've started exploring SQL capabilities, beginning with the fundamental *SELECT FROM WHERE* incantation
  - ▶ Focused on the **relational algebra semantics** of the underlying “selection”, “projection”, and “Cartesian product” operators
- ▶ Today
  - ▶ More “basic SQL” material, especially SQL **set operations**
  - ▶ Semantics and implications of NULL
- ▶ But first, short, but important “miscellania”
  - ▶ The material isn't very deep ... (“syntax”, not “semantics”)
  - ▶ These are “pointer slides”: be aware of the material, read textbook & Internet as necessary



# Ambiguous Attribute Names

- ▶ Different relations may use the same attribute name in their schema
  - ▶ That's perfectly legal ☺
- ▶ **Q:** but what if your query must refer to both relations in the same query and must distinguish between the identical attribute name?
- ▶ **A:** resolve syntactic **ambiguity** by qualifying the attribute name with the appropriate relation name

```
1 SELECT Fname, EMPLOYEE.Name, Address
2 FROM EMPLOYEE, DEPARTMENT
3 WHERE DEPARTMENT.Name = 'Research' AND DEPARTMENT.Dnumber
      = EMPLOYEE.Dnumber;
```

## Aliasing and Renaming (I)

- ▶ The ambiguity issue can even arise with respect to relation names
- ▶ Example: *“For each employee, retrieve the employee’s first and last name and the first and last name of his or her immediate supervisor”*
  - ▶ The problem: the query refers to the same relation twice
- ▶ Solution: have the query declare alternative relation names
  - ▶ These are called aliases or tuple variables
  - ▶ Syntax: use the AS clause
  - ▶ (Note: the AS may be dropped in most SQL implementations, but don’t do that! ☺)

```
1 SELECT E.Fname , E.Lname , S.Fname , S.Lname
2    FROM EMPLOYEE AS E, EMPLOYEE AS S
3    WHERE E.Super_ssn = S.Ssn;
```

## Aliasing and Renaming (II)


- ▶ You can use this renaming mechanism in any query to specify tuple variables for every table in the WHERE clause
  - ▶ Even when there no ambiguity exists 😊
  - ▶ Recommended practice to make your queries more readable
- ▶ Example: *“Retrieve the name and address of all employees who work for the ‘Research’ department”*

```
1 SELECT E.Fname, E.LName, E.Address
2 FROM EMPLOYEE AS E, DEPARTMENT AS D
3 WHERE D.DName = 'Research' AND D.Dnumber = E.Dno;
```

## More selection examples

```
SELECT *  
FROM Product  
WHERE prod_price IS NOT NULL;
```

```
SELECT *  
FROM Product  
WHERE prod_price BETWEEN 20 AND 40;
```



Inclusive

```
SELECT *  
FROM Product  
WHERE prod_price > 20 AND prod_manufacturer = 'GizmoWorks';
```

```
SELECT *  
FROM Product  
WHERE prod_manufacturer IN ('GizmoWorks', 'WidgetsRUs');
```



## Selection with pattern matching

```
SELECT *  
FROM Product  
WHERE prod_name LIKE '%Gizmo%';
```

- % = Match any sequence of 0-or-more characters
- \_ = Match any single character
- [abc] = Match any one character listed
- [a-c] = Match any one character in range

# Making results distinct

```
SELECT DISTINCT manufacturer  
FROM Product;
```

## Product

prod_name	prod_price	prod_manufacturer
Gizmo	\$19.99	GizmoWorks
Powergizmo	\$39.99	GizmoWorks
Widget	\$19.99	WidgetsRUs
HyperWidget	\$203.99	Hyper



prod_manufacturer
GizmoWorks
WidgetsRUs
Hyper

Ensures results are a set.

## Computation in SELECT clauses

```
SELECT location, time, celsius * 1.8 + 32 AS fahrenheit  
FROM SensorReading;
```

Use AS. Otherwise, get  
a default field name.

```
SELECT player_id, Floor(height) AS feet, (height - Floor(height)) * 12 AS inches  
FROM Player;
```

```
SELECT student_id, CASE WHEN lastname < 'N' THEN 1 ELSE 2 END AS group  
FROM Student;
```

```
SELECT lastname || ', ' || firstname AS name  
FROM Student;
```

## Syntax details

- Strings use single quotes, not double
- Equality test uses single =, not double
- Amount of whitespace doesn't matter

'Houston'

x = 5

# Sorting

```
SELECT prod_name, prod_manufacturer  
FROM Product  
ORDER BY prod_price DESC, prod_manufacturer;
```

## Product

prod_name	prod_price	prod_manufacturer
Gizmo	\$19.99	GizmoWorks
Powergizmo	\$39.99	GizmoWorks
Widget	\$19.99	WidgetsRUs
HyperWidget	\$203.99	Hyper



prod_name	prod_manufacturer
HyperWidget	Hyper
Powergizmo	GizmoWorks
Gizmo	GizmoWorks
Widget	WidgetsRUs

## Sorting on computation results

```
SELECT item, price * quantity AS total  
FROM Order  
ORDER BY price * quantity;
```

```
SELECT item, price * quantity AS total  
FROM Order  
ORDER BY total;
```

### Order

item	price	quantity
apple	\$0.50	3
orange	\$0.60	2
banana	\$0.40	4
peach	\$0.80	1



item	total
peach	\$0.80
orange	\$1.20
apple	\$1.50
banana	\$1.60

# Multiset semantics vs. Sorting

Table rows are unordered, except when they're ordered.



More accurately, unless you use ORDER BY:

- Can't assume anything about ordering.
- Ordering depends on implementation, which can vary.
- Query results don't necessarily maintain order of original table.

## Subset of results

### Product

prod_name	prod_price	prod_manufacturer
Gizmo	\$19.99	GizmoWorks
Powergizmo	\$39.99	GizmoWorks
Widget	\$19.99	WidgetsRUs
HyperWidget	\$203.99	Hyper

```
SELECT prod_name, prod_manufacturer  
FROM Product  
LIMIT 2;
```

Some SQLs:  
SELECT TOP 2 ...

```
SELECT prod_name, prod_manufacturer  
FROM Product  
ORDER BY prod_price  
LIMIT 2;
```

prod_name	prod_manufacturer
Gizmo	GizmoWorks
HyperWidget	Hyper

Arbitrary  
products

prod_name	prod_manufacturer
Gizmo	GizmoWorks
Widget	WidgetsRUs

Cheapest  
products



# String Matching

- ▶ You can always do an **exact match** on string value

```
1      delete from instructor
2      where dept_name='Finance'
```

- ▶ SQL includes a **string-matching operator** for comparisons on character strings
- ▶ The LIKE operator uses patterns that are described using two special characters
  - ▶ Percent (%): matches any **substring**

```
1  select name
2  from instructor where name like '%dar%'
```

- ▶ Can use the ESCAPE operator to define the “escape” character and thus allow comparison against the % character itself

```
1  like '100\%' escape '\'
```

- ▶ Underscore (\_): matches any **character**

## String Comparisons: Some Examples

- ▶ Even though SQL in general is **case insensitive**, patterns are **case sensitive**!
- ▶ Pattern matching examples
  - ▶ 'Intro%' matches any string beginning with "Intro"
  - ▶ '%Comp%' matches any string containing "Comp" as a substring
  - ▶ '\_\_\_' matches any string of exactly three characters
  - ▶ '\_\_\_%' matches any string of at least three characters
- ▶ Many built-in String functions
  - ▶ Concatenation
  - ▶ Convert between "upper case" and "lower case"
  - ▶ Substring extraction
  - ▶ String length



# Introduction

- ▶ You may have heard that SQL is based on the relational algebra which, in turn, is based on a formal **set-based** definition of relations and tuples ☺
- ▶ We've already discussed the “selection” and “projection” operators which were developed specifically for relational databases
- ▶ We've already mentioned that the semantics of “Cartesian product” are changed from classic set theory
  - ▶ The Cartesian product of a set of  $n$ -tuples with a set of  $m$ -tuples yields a set of “**flattened**”  $(n + m)$ -tuples
  - ▶ “Basic” set theory would require a set of 2-tuples, each containing an  $n$ -tuple and an  $m$ -tuple
- ▶ Unsurprisingly, SQL also includes the set **union**, set **difference**, and set **intersection** operators
  - ▶ Respectively: SQL UNION, EXCEPT (some dialects have MINUS), and INTERSECT
- ▶ However: SQL adds an additional constraint to these operators: **union compatibility**

## Union Compatibility

- ▶ In order to be used in a UNION, the two relations must have the “same attribute characteristics”
  - ▶ The attributes and their domains must be compatible: they **share the same number of columns** and their corresponding columns **share the same or compatible domains**
- ▶ This property holds for the INTERSECT and EXCEPT operators as well

### Examples:

- ▶  $R(id, name)$  and  $S(id, name, grades)$  are not union compatible:  $R$  has 2 attributes, but  $S$  has 3 attributes
- ▶  $R(id, name)$  and  $S(id, grades)$  are not union compatible: the domains of “name” and “grades” are different (“string” versus “numeric”)
- ▶  $R(id, name)$  and  $S(id, StudentName)$  are union compatible:  $R$  and  $S$  both have 2 attributes and their domains are also identical (only the column names differ)

# Set Operations: Examples

*Find courses that ran in Fall 2009 **or** in Spring 2010*

```
1      (select course_id from section
2      where sem = 'Fall' and year = 2009)
3      union
4      (select course_id from section
5      where sem = 'Spring' and year = 2010)
```

*Find courses that ran in Fall 2009 **and** in Spring 2010*

```
1      (select course_id from section
2      where sem = 'Fall' and year = 2009)
3      intersect
4      (select course_id from section
5      where sem = 'Spring' and year = 2010)
```

*Find courses that ran in Fall 2009 **but not** in Spring 2010*

```
1      (select course_id from section
2      where sem = 'Fall' and year = 2009)
3      except
4      (select course_id from section
5      where sem = 'Spring' and year = 2010)
```

## Tables as Sets in SQL (I)

- ▶ The “multi-set” issue is one of the most important differences between “commercial” SQL and the formal relational algebra
- ▶ This course eschews “DML formalism” as much as possible, and therefore focuses on SQL’s way of doing things ...
  - ▶ But you must be aware that SQL allows a relation to include multiple tuples that are identical in all their attribute values
- ▶ This implies that an SQL table is not a set of tuples
  - ▶ Instead: an SQL table is a multi-set (or “bag”) of tuples
  - ▶ By default, SQL will not eliminate duplicate tuples
- ▶ Note: we can constrain some SQL relations to be sets because of a “key constraint” or because we use DISTINCT in a SELECT statement

## Tables as Sets in SQL (I)

- ▶ Q: why doesn't SQL simply eliminate duplicates in result-sets?
- ▶ A<sub>1</sub>: because sometimes the client does want to see duplicate tuples 😊
  - ▶ Example: you do a JOIN between a “users” and “tasks” tables, and want to get all the tasks associated with a given user
  - ▶ See discussion [here](#)
  - ▶ ОТОН: sometimes the client is introducing a “join duplication” bug 😞
- ▶ A<sub>2</sub>: because that can be a **very expensive operation** 😞
  - ▶ Consider implementation: e.g., “first sort, then eliminate duplicates”
- ▶ Hold this discussion in mind when we (subsequent lecture) discuss **aggregate functions**
  - ▶ Typically we don't want to eliminate duplicates



## Set Operations & Issue of Duplicates

- ▶ **UNION** operation automatically eliminates duplicates from the result relation
  - ▶ To override this behavior: must specify UNION ALL
- ▶ **INTERSECT** operation automatically eliminates duplicates from the result relation
  - ▶ To override this behavior: must specify INTERSECT ALL
  - ▶ The number of duplicate tuples is then the minimum of the number of duplicates in the two relations
- ▶ **EXCEPT** operation automatically eliminates duplicates from the result relation
  - ▶ To override this behavior: must specify EXCEPT ALL
  - ▶ The number of duplicate tuples is then the number of duplicates in the first relation minus the number of duplicates in the second relation
    - ▶ So long as that difference is **positive** 😊

# Tables are *multisets*

So, what do we know about multisets?

## Two ways to think about multisets

Tuple
(1, a)
(1, a)
(1, b)
(2, c)
(2, c)
(2, c)
(1, d)
(1, d)



Tuple	$\lambda(X)$
(1, a)	2
(1, b)	1
(2, c)	3
(1, d)	2

# Multiset union

Tuple	$\lambda(X)$
(1, a)	2
(1, b)	0
(2, c)	3
(1, d)	0

U

Tuple	$\lambda(Y)$
(1, a)	5
(1, b)	1
(2, c)	2
(1, d)	2

=

Tuple	$\lambda(Z)$
(1, a)	7
(1, b)	1
(2, c)	5
(1, d)	2

$$\lambda(Z) = \lambda(X) + \lambda(Y)$$

# Multiset intersection

Tuple	$\lambda(X)$
(1, a)	2
(1, b)	0
(2, c)	3
(1, d)	0

$\cap$

Tuple	$\lambda(Y)$
(1, a)	5
(1, b)	1
(2, c)	2
(1, d)	2

$=$

Tuple	$\lambda(Z)$
(1, a)	2
(1, b)	0
(2, c)	2
(1, d)	0

$$\lambda(Z) = \min(\lambda(X), \lambda(Y))$$

# Multiset difference

Tuple	$\lambda(X)$
(1, a)	2
(1, b)	0
(2, c)	3
(1, d)	0

—

Tuple	$\lambda(Y)$
(1, a)	5
(1, b)	1
(2, c)	2
(1, d)	2

=

Tuple	$\lambda(Z)$
(1, a)	0
(1, b)	0
(2, c)	1
(1, d)	0

$$\lambda(Z) = \lambda(X) - \lambda(Y)$$

## SQL syntax

Sets:

```
SELECT a  
FROM R  
UNION  
SELECT a  
FROM S;
```

```
...  
INTERSECT  
...
```

```
...  
EXCEPT  
...
```

Multisets:

```
SELECT a  
FROM R  
UNION ALL  
SELECT a  
FROM S;
```

```
...  
INTERSECT ALL  
...
```

```
...  
EXCEPT ALL  
...
```

## Segue: The Problem Of NULL

- ▶ Remainder of lecture is our “drill-down” into the topic of NULL
- ▶ Besides its intrinsic importance ... this topic is connected to our discussion of SQL set operations
  - ▶ As we shall see, the semantics of NULL imply that the semantics of “duplicates” is much more complicated than we’d like ☹



# The Problem of NULL

# Introduction

*The simple scientific fact is that an SQL table that contains a null isn't a relation; thus, relational theory doesn't apply, and all bets are off.*

*Chris Date*

## NULL in a nutshell

- ▶ Each domain is augmented with a NULL
- ▶ NULL, intuitively stands for one of the following
  - ▶ Value **unknown**
  - ▶ Value **not permitted to be known** (to some of us)
  - ▶ Value **not applicable**

We've been using NULL in many of our examples: what is the fuss all about?

# The Problem of NULL: What Semantics Is It Modeling?

- ▶ We just said that “NULL can have several interpretations”
  - ▶ Example: a query returns NULL as the value of the `hair_color` attribute
- ▶ Does that mean that the person is bald?
- ▶ Or: the person has hair, but you just don't know what color?
- ▶ Or: can the person be bald or have hair, but you just don't know which one applies?
- ▶ Or: the person is in the midst of a hair coloring exercise and you **only temporarily** don't know the color?
- ▶ Or: (even more likely), the data-collection person simply forgot to record the data 😊

“Hair color” example taken from [here](#)

# Dealing With The Consequences Of NULL

- ▶ SQL is forced to replace traditional **boolean** logic with a new **3-Valued** logic
  - ▶ Once we introduce NULL (regardless of whether this is a good idea or not) ...
  - ▶ We have no choice but to change very basic ideas of how logic behaves ☹
- ▶ Abbreviations
  - ▶ **T** for TRUE
  - ▶ **F** for FALSE
  - ▶ **U** for UNKNOWN
- ▶ Let's start with a **boolean logic** refresher ☺

	NOT		OR	F	T		AND	F	T
F	T	F	F	F	T	F	F	F	F
T	F	T	T	T	T	T	F	T	T

## Three-Valued Logic: Intuition

- ▶ May help your intuition if you think of **U** as being “in between” **F** and **T**
  - ▶ But more accurate interpretation on next slide
- ▶ If you think of
  - ▶  $\text{NOT}(x)$  as  $1 - x$
  - ▶  $x \text{ OR } y$  as  $\max(x, y)$
  - ▶  $x \text{ AND } y$  as  $\min(x, y)$
- ▶ For 2-valued logic we have
  - ▶ FALSE is 0
  - ▶ TRUE is 1
- ▶ For 3-valued logic we have
  - ▶ FALSE is 0
  - ▶ **UNKNOWN** is 0.5
  - ▶ TRUE is 1

## Three-Valued Logic

Instead of thinking of NULL as being “in between” F and T, better to say “NULL means: maybe T or maybe F”

	NOT
F	T
U	U
T	F

	OR	F	U	T
F	F	F	U	T
U	U	U	U	T
T	T	T	T	T

	AND	F	U	T
F	F	F	F	F
U	F	F	U	U
T	F	F	U	T

## NULL & SELECT-FROM-WHERE statement

- ▶ Each tuple in the FROM clause is tested against the WHERE predicate
- ▶ Here are the rules
  - ▶ If  $P(\text{tuple})$  is TRUE, it is passed to SELECT
  - ▶ If  $P(\text{tuple})$  is UNKNOWN, it **is not passed** to SELECT
  - ▶ If  $P(\text{tuple})$  is FALSE, it **is not passed** to SELECT
- ▶ Summary: For SELECT queries, UNKNOWN behaves exactly the same as FALSE
- ▶ Unfortunately this behavior **is different** for DDL and INSERT statements
  - ▶ Where UNKNOWN behaves as TRUE ☹

## NULL & SELECT: AND Example

Any comparison in which one side is NULL is UNKNOWN

R	A	B	C
1	6	8	
2	7	9	
3	NULL	8	
4	NULL	9	

```
select A from R where B = 6 AND C = 8
```

	A
	1



## NULL & SELECT: OR Example

Any comparison in which one side is NULL is UNKNOWN

R	A	B	C
	1	6	8
	2	7	9
	3	NULL	8
	4	NULL	9

```
select A from R where B = 6 OR C = 8
```

	A
	1
	3

## NULL & SELECT: NULL = NULL is FALSE!

Any comparison in which one side is NULL is UNKNOWN

R	A	B	C
	1	6	8
	2	7	9
	3	NULL	8
	4	NULL	9

```
select A from R where B = NULL;
```

The result of this query is: **empty table!**


## NULL & SELECT: NULL $\neq$ NULL is FALSE!

Any comparison in which one side is NULL is UNKNOWN

R	A	B	C
	1	6	8
	2	7	9
	3	NULL	8
	4	NULL	9

```
select A from R where B <> NULL;
```

The result of this query is: **empty table!**



## NULL & SELECT: NULL Doesn't Equal Itself!

Any comparison in which one side is NULL is UNKNOWN

R	A	B	C
1	6	8	
2	7	9	
3	NULL	8	
4	NULL	9	

```
select A from R where B = B;
```

Because, going row by row:

1. 6 = 6 is TRUE
2. 7 = 7 is TRUE
3. NULL = NULL is UNKNOWN
4. NULL = NULL is UNKNOWN

	A
1	
2	

## Alleviate Some of the Problem: Introduce New Keywords

- ▶ A new keyword made of three words: **IS NOT NULL**
- ▶ A new keyword made of two words: **IS NULL**

## IS NOT NULL Example

R	A	B	C
	1	6	8
	2	7	9
	3	NULL	8
	4	NULL	9

```
select A from R where B is not null
```

	A
	1
	2

# IS NULL Example

R	A	B	C
1	6	8	
2	7	9	
3	NULL	8	
4	NULL	9	

```
select A from R where B is null
```

A
3
4

## NULL & Arithmetic Wierdness

If one of the operands is NULL, the result is NULL!

- ▶  $5 + \text{NULL} = \text{NULL}$
- ▶  $\text{NULL} - \text{NULL} = \text{NULL}$
- ▶  $0 * \text{NULL} = \text{NULL}$
- ▶  $\text{NULL} / 0 = \text{NULL}$

Key point: given the semantics of NULL, these results are inevitable



# NULL & Duplicates

## Quick review of SQL & duplicates

- ▶ Standard SELECT FROM WHERE statement **does not remove duplicates** at any stage of its execution
  - ▶ Use SELECT **DISTINCT** FROM WHERE to remove duplicates
- ▶ Standard UNION, EXCEPT, INTERSECT **do remove duplicates**
- ▶ UNION ALL, EXCEPT ALL, INTERSECT ALL **do not remove duplicates** with rather interesting semantics
- ▶ All NULLs are **duplicates of one another**
  - ▶ Implication: so if you use DISTINCT, you'll get at most NULL tuple
  - ▶ Even though it is UNKNOWN whether they are equal to each other 😊

# Today's Lecture: Wrapping it Up

Miscellanea

Set Operations

The Problem of NULL

- ▶ “Miscellania” material mostly covered in Textbook, Chapter 3.4
- ▶ Set operations covered in Textbook, Chapter 3.5
- ▶ NULL covered in Textbook, Chapter 3.6