

# Transaction Isolation: Two-Phase Locking & Deadlocks

COM 3563: Database Implementation


Avraham Leff

Yeshiva University

*avraham.leff@yu.edu*

COM3563: Fall 2020





- ▶ Previous lecture: we began a series of lectures on “how to implement transactions”
- ▶ Presented the concept of **serializable schedules**: schedules that are equivalent to some **serial schedule**
- ▶ We proceeded to define “equivalence” in terms of **conflict serializability**
  - ▶ “Can the concurrent schedule *be transformed into a serial schedule* after swapping of non-conflicting operations?”
- ▶ We provided a test to verify conflict serializability: “is the schedule’s **dependency graph** acyclic?”

# Today

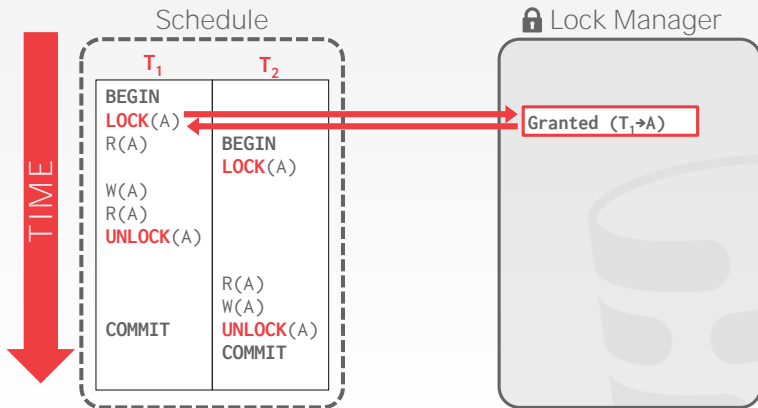
- ▶ Previous lecture provided us with solid theoretical foundation for concurrency control 😊
- ▶ But: no practical guidance for how the DBMS should construct a schedule 😞
- ▶ After all: testing a schedule for serializability **after it has executed** (or even **after it's been constructed**) is a little too late!
- ▶ Our goal is to develop concurrency control protocols that will automatically guarantee serializability
  - ▶ Without needing to inspect a dependency graph



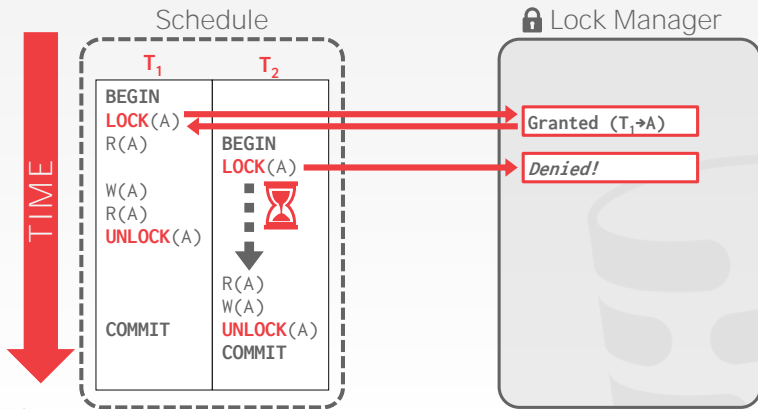
## Locks: The Basic Idea

- ▶ Use locks to protect database objects from **concurrent access**
- ▶ This is a “pessimistic” approach to concurrency control: ensure that the DBMS doesn’t even create a schedule containing **conflicts**
- ▶ Think of a lock as a variable that’s associated with a data item
- ▶ The lock describes the status of that data item with respect to “*what operations can be applied now?*”
- ▶ We’ll begin by assuming that locks have **binary state**: either “locked” or “unlocked”
  - ▶ Locked ⇒ item cannot be accessed by another thread
  - ▶ Unlocked ⇒ item can be accessed
  - ▶ We’ll change that assumption very soon 😊
- ▶ We’ll begin by assuming “one lock per data item”
  - ▶ Relax that assumption later

# EXECUTING WITH LOCKS

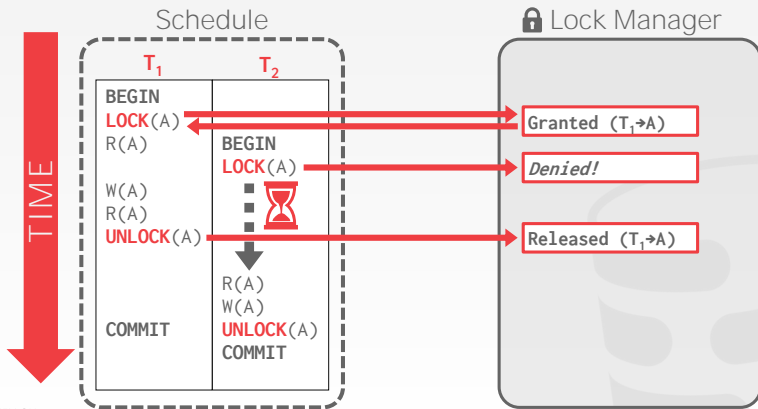


# EXECUTING WITH LOCKS

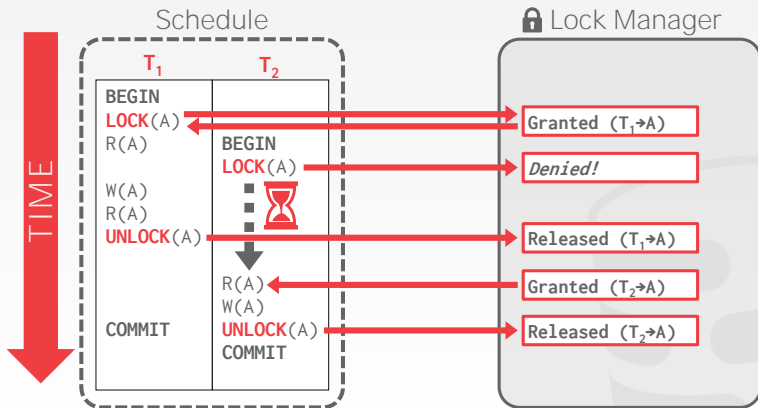




# EXECUTING WITH LOCKS



# EXECUTING WITH LOCKS



## Division of responsibilities:

- ▶ Transactions request locks (or “lock upgrades”)
- ▶ Lock manager responds to requests from transaction: either **grants** or **blocks** requests
- ▶ Transactions responsible for **releasing locks** when no longer needed
- ▶ Lock manager maintains an internal **lock-table**
  - ▶ Updates the lock table in response to lock request/release API invocations
- ▶ The lock-table tracks *“which transactions hold which locks?”*
- ▶ The lock table also tracks *“which transactions are waiting to acquire locks?”*
  - ▶ We refer to a tx that’s “waiting” as a **“blocked transaction”**

## Enrich Lock Semantics With “Read Versus Write”

- ▶ We previously said that a lock has “binary state”
- ▶ In practice, that model is too restrictive for DBMS use
  - ▶ After all: why not allow **multiple readers** to access the same data concurrently?
- ▶ So: enrich model to allow data items to be locked in one of **two modes**
  - ▶ **Exclusive (X) mode**: data item can be both read as well as written
  - ▶ **Shared (S) mode**: data item can only be read
- ▶ Now the question of whether a lock **should or should not be granted** depends on
  - ▶ *“Is the data item already locked?”*
  - ▶ *“In what mode is the datum locked?”*

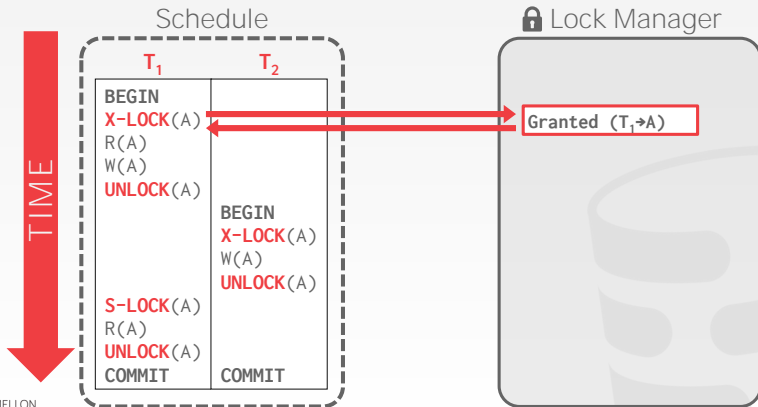
Unlike non-transactional models, programmers do **not make explicit lock requests to the DBMS**: these are generated automatically 😊

## Lock Compatibility Matrix

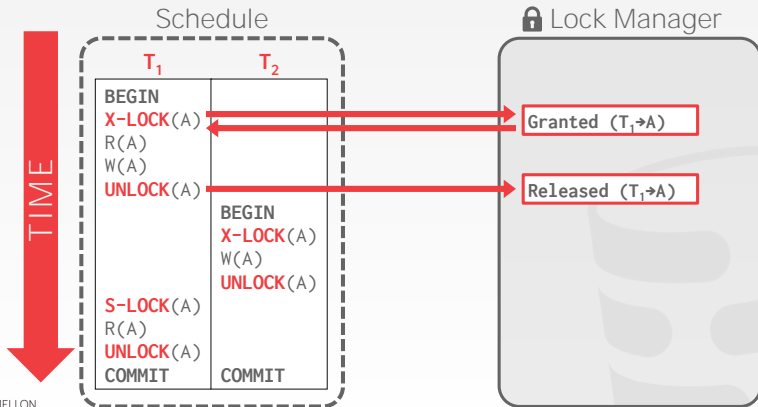
	S	X
S	true	false
X	false	false

- ▶ Transaction is granted a lock *iff* requested lock is **compatible** (see Figure above) with locks already held on the item by other transactions
- ▶ Any number of transactions can hold **shared locks** on an item
- ▶ One transaction holds an exclusive lock on the item → **no other transaction will be granted a lock**
- ▶ Requesting transaction blocks if lock is not granted
  - ▶ System tracks status of existing, incompatible, locks held by other transactions
  - ▶ When all other locks have been released, transaction is granted lock, and resumes execution

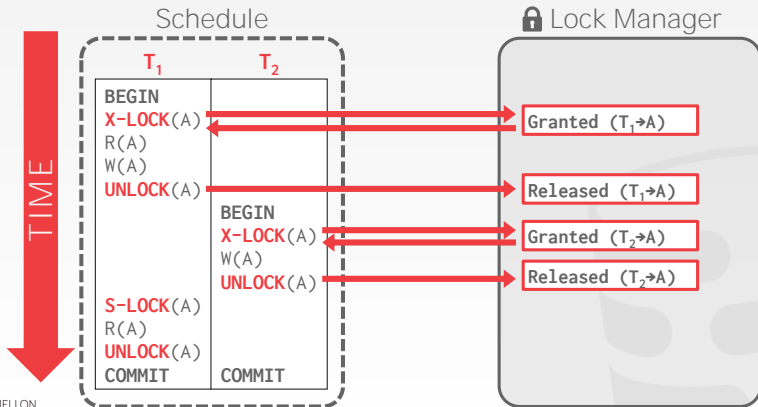
# EXECUTING WITH LOCKS



# EXECUTING WITH LOCKS

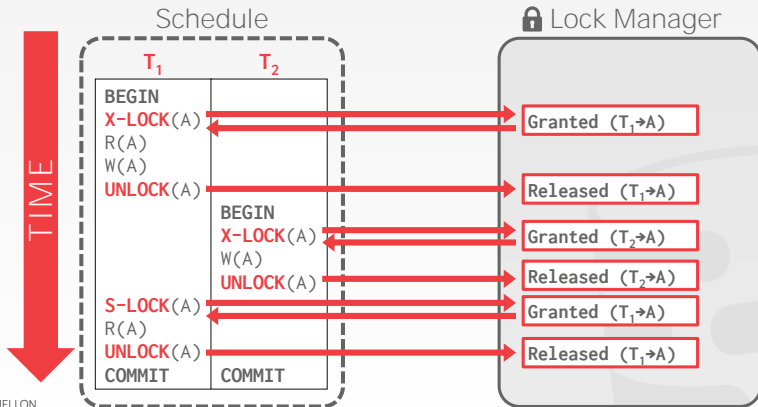


# EXECUTING WITH LOCKS





# EXECUTING WITH LOCKS



## Locking (By Itself) Is Not A CC Protocol!


```
 $T_2$ : lock-S(A);  
      read (A);  
      unlock(A);  
      lock-S(B);  
      read (B);  
      unlock(B);  
      display(A+B)
```

- ▶ Important: locking as performed in the Figure above **does not guarantee serializability!**
- ▶ Example: if  $A$  and  $B$  are **written** by another tx in-between **read** of  $A$  and  $B$ , displayed sum would be wrong
- ▶ We need to use a **locking protocol** to further restrict the set of possible schedules
  - ▶ “A set of rules followed by all transactions while requesting and releasing locks”
- ▶ Goal: a locking protocol that **guarantees conflict serializability**

## Two-Phase Locking Protocol

- ▶ The **two-phase locking protocol** does guarantee conflict-serializable schedules
- ▶ **Phase 1:** The “growing phase”
  - ▶ Transaction **may obtain** locks
  - ▶ Transaction **may not release** locks
- ▶ **Phase 2:** The “shrinking phase”
  - ▶ Transaction **may release** locks
  - ▶ Transaction **may not obtain** locks
- ▶ This version is “vanilla” 2PL: allows transactions to release locks **even before commit point**
- ▶ The protocol assures serializability!
- ▶ Can be proven that the transactions can be serialized in the **order of their lock points**
  - ▶ The point where a transaction acquired its final lock
  - ▶ **Important:** “lock point”  $\neq$  “commit point” (hold that thought)

# EXECUTING WITH LOCKS



## Lock Manager

Granted ( $T_1 \rightarrow A$ )

Released ( $T_1 \rightarrow A$ )

Granted ( $T_2 \rightarrow A$ )


Released ( $T_2 \rightarrow A$ )

Granted ( $T_1 \rightarrow A$ )

Released ( $T_1 \rightarrow A$ )


## TWO-PHASE LOCKING

The txn is not allowed to acquire/upgrade locks after the growing phase finishes.

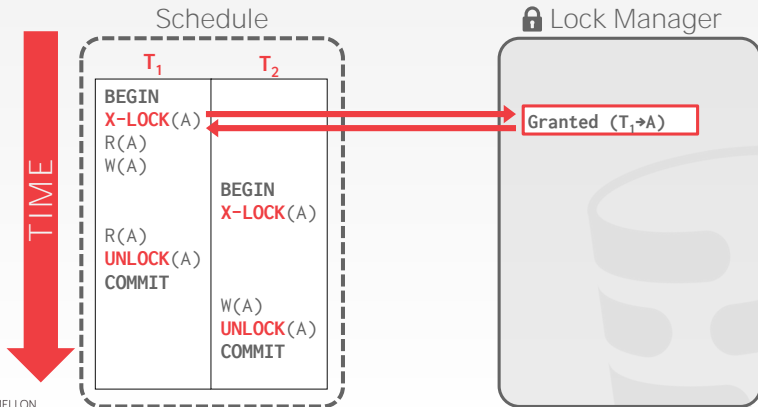


## TWO-PHASE LOCKING

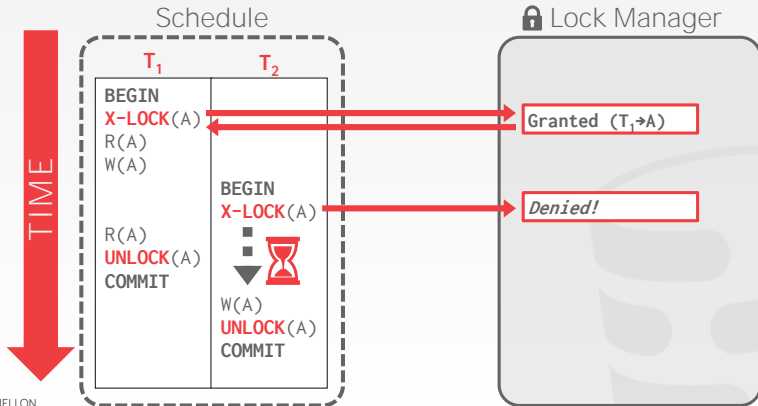
The txn is not allowed to acquire/upgrade locks after the growing phase finishes.



## EXECUTING WITH 2PL

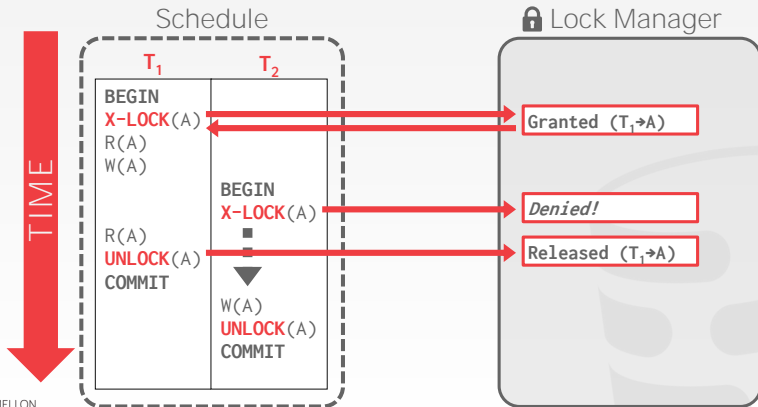


# EXECUTING WITH 2PL

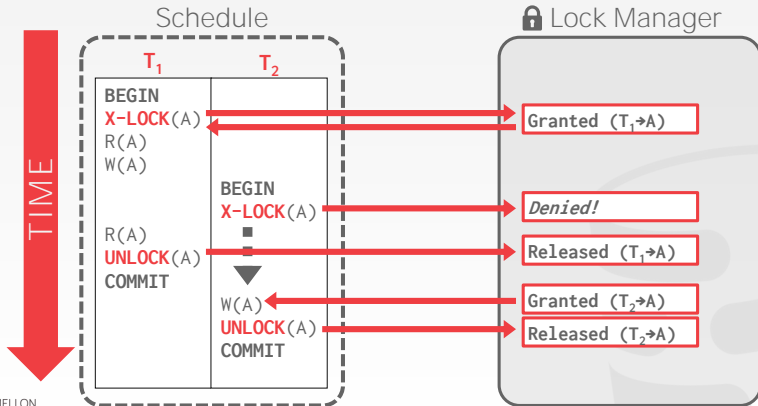




# EXECUTING WITH 2PL




# EXECUTING WITH 2PL



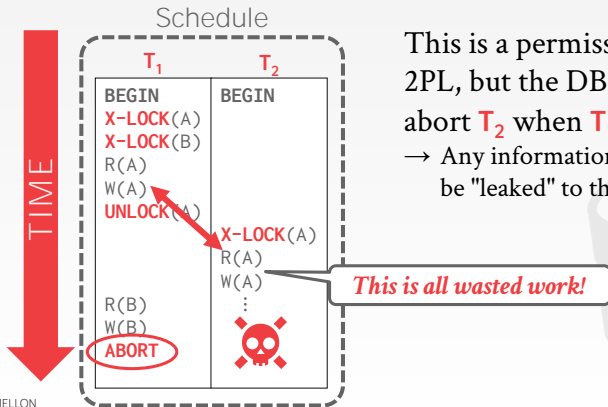
## 2PL: The Good News & The Bad

- ▶ If every transaction in a schedule follows the two-phase locking protocol, then the resulting schedule is guaranteed to be serializable
  - ▶ We can show that 2PL generates schedules whose dependency graph is **acyclic**
  - ▶ 2PL may limit the amount of concurrency that can occur in a schedule because some serializable schedules will be prohibited by two-phase locking protocol
- ▶ We can prove that txs in a 2PL schedule are serialized in the **order of their lock points**
  - ▶ Meaning: the point where a transaction acquired its final lock
- ▶ Unfortunately: 2PL schedules are subject to **cascading aborts** (last lecture)
  - ▶ The problem: DBMS doesn't prevent other txs from seeing the **"dirty data"** affected by the aborted tx

## 2PL – CASCADING ABORTS



## 2PL – CASCADING ABORTS



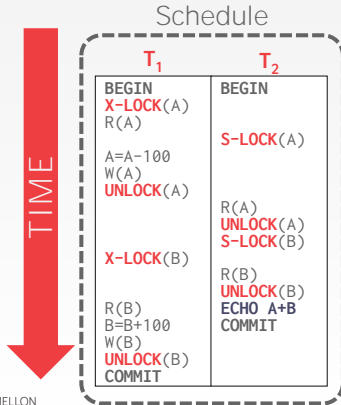
This is a permissible schedule in 2PL, but the DBMS has to also abort  $T_2$  when  $T_1$  aborts.

→ Any information about  $T_1$  cannot be "leaked" to the outside world.

## Variations On 2PL

- ▶ **Conservative 2PL** (or “static”) requires that a tx lock all the items it accesses before the tx begins
  - ▶ This variant requires that the tx “pre-declare” its read-set and write-set
  - ▶ Advantage: this variant is a **deadlock-free** protocol
  - ▶ (More on deadlocks later)
- ▶ **Strict 2PL**: tx does not release **exclusive locks** until after it commits or aborts
  - ▶ Advantage: doesn't incur **cascading aborts** and ensures “**recoverability**”
  - ▶ Advantage: aborted txs can be undone simply by **restoring original values of modified data**
- ▶ **Rigorous 2PL**: tx holds all locks until completion
  - ▶ Txs can be serialized in the order in which they commit
- ▶ Note: most databases implement **rigorous 2PL**, but refer to it as simply “two-phase locking”

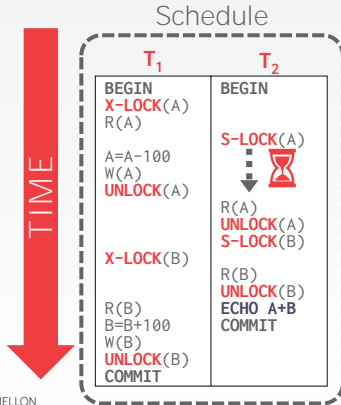
# NON-2PL EXAMPLE



Initial Database State

A=1000, B=1000

# NON-2PL EXAMPLE

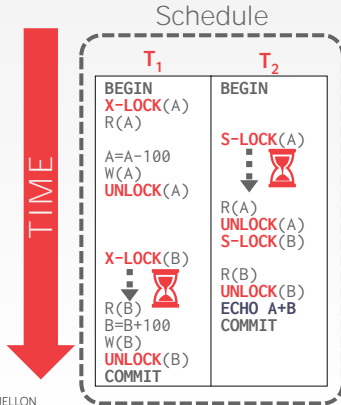


Initial Database State

**A**=1000, **B**=1000




# NON-2PL EXAMPLE



Initial Database State

**A**=1000, **B**=1000

# NON-2PL EXAMPLE



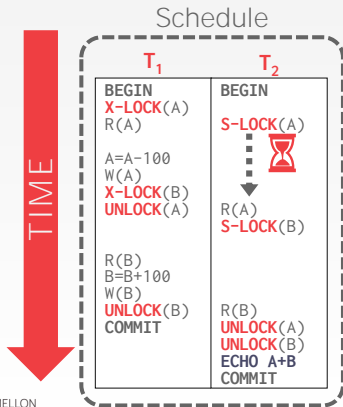
Initial Database State

A=1000, B=1000

$T_2$  Output

A+B=1100


## 2PL EXAMPLE



Initial Database State

**A**=1000, **B**=1000


## 2PL EXAMPLE



Initial Database State

**A**=1000, **B**=1000

## 2PL EXAMPLE



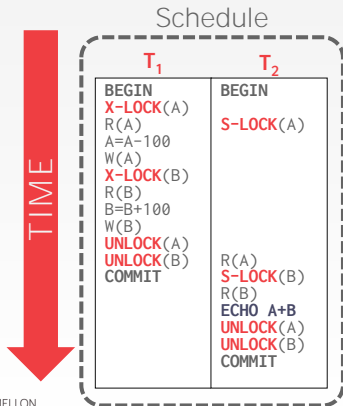
Initial Database State

**A=1000, B=1000**

$T_2$  Output

**A+B=2000**


# STRICT 2PL EXAMPLE



Initial Database State

A=1000, B=1000


# STRICT 2PL EXAMPLE



Initial Database State

A=1000, B=1000

# STRICT 2PL EXAMPLE



Initial Database State

**A**=1000, **B**=1000

$T_2$  Output

**A+B**=2000





## Deadlocks: A Real Problem

- ▶ A **deadlock** is a cycle of transactions waiting for locks to be released by each other
  - ▶ Trivially triggered by schedules that don't require txs to "pre-declare" their read-sets and write-sets
- ▶ Two ways of dealing with deadlocks
  - ▶ Deadlock **detection**: DBMS allows deadlocks to occur, "terminates with extreme prejudice" when it detects deadlock situation
  - ▶ Deadlock **prevention**: when a tx tries to acquire a lock that is currently held by another tx, DBMS kills one of them to **prevent a deadlock**

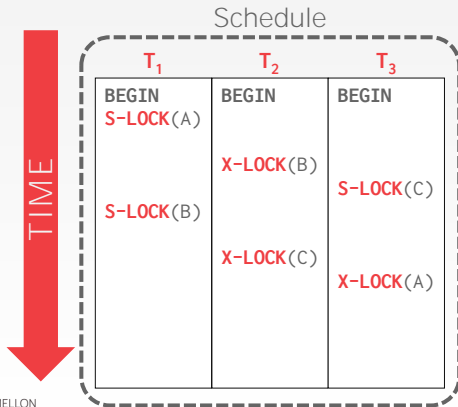
## Deadlock Detection: “Waits-For” Graphs

- ▶ DBMS creates a “waits-for” graph to keep track of what locks each txn is waiting to acquire
  - ▶ Graph nodes are transactions
  - ▶ Graph inserts an edge from  $T_i \Rightarrow T_j$  “iff  $T_i$  is waiting for  $T_j$  to release a lock”
- ▶ DBMS periodically check for cycles in its “waits-for” graph and then make a decision on how to break the deadlock (next slide)
- ▶ The decision: select a “victim” txn to rollback that **breaks the cycle** in the “waits-for” graph
  - ▶ Victim is either restarted afterwards or simply aborted
- ▶ Note the trade-off between how often the DBMS invokes the “deadlock detection” algorithm and how long txs **have to wait** while they’re deadlocked


## Deadlock Handling: “Victim” Selection

- ▶ The algorithm has many variables it can tune (or ignore)
- ▶ Examples: select the victim ...
  - ▶ “By age” (most recent tx)
  - ▶ “By progress” (tx with fewest operations performed so far)
  - ▶ “By # of items already locked”
  - ▶ “By the # of txs that the DBMS will have to rollback with it”
- ▶ In addition to optimizing for “overall performance”, the algorithm may choose “fairness”
  - ▶ Example: *“the number of times that the tx has previously been aborted due to deadlocks”*
  - ▶ Systems that don’t consider this factor may suffer from “tx starvation”

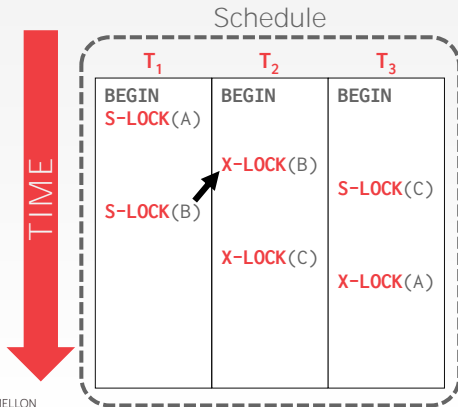
# DEADLOCK DETECTION




Waits-For Graph




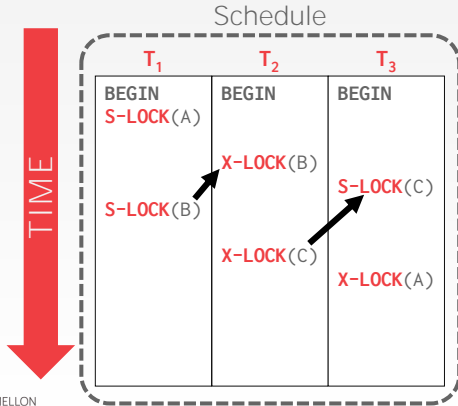
# DEADLOCK DETECTION




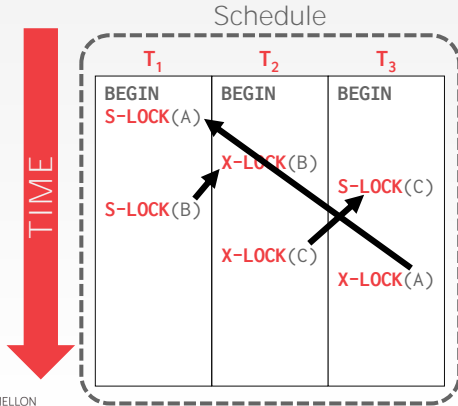
Waits-For Graph



# DEADLOCK DETECTION



# DEADLOCK DETECTION





## Deadlock Prevention

- ▶ No need to **detect** a deadlock: as soon as  $tx_i$  requests a lock that is currently held by  $tx_j$ , DBMS aborts one of these txs
- ▶ But which tx should be aborted?
- ▶ Can assign priorities based on timestamps (“older tx has **higher priority**”)
  - ▶ Then based on such priorities, implement one of two algorithms
- ▶ Wait-Die (or “Old Waits for Young”)
  - ▶ If *requesting tx* has higher priority than *holding tx*, then *requesting tx* **waits for holding tx**
  - ▶ Otherwise: younger requesting tx **aborts**
- ▶ Wound-Wait (or “Young Waits for Old”)
  - ▶ If *requesting tx* has higher priority than *holding tx*, then *holding tx* **aborts** (and releases lock)
  - ▶ Otherwise: younger requesting tx **waits**

## Protocol Similarities

- ▶ In both **wait-die** and **wound-wait** schemes, the rolled back tx is **restarted with its original timestamp**
  - ▶ Eventually, the aborted (younger) transactions will become the oldest transactions in the system and complete successfully
  - ▶ So: both protocols are therefore **fair**
- ▶ In both schemes **older txs “beat” younger txs**
- ▶ **Q:** can you suggest why we favor older txs this way?
- ▶ **A:** older txs (having run longer) typically hold more locks and have read/written more data than younger txs
  - ▶ So relatively more expensive to rollback older txs

Both deadlock prevention schemes are based on the concept of **imposing a partial ordering of all data items**, then requiring that txs can only lock data items **in the order specified by the partial order**

## Protocol Differences


- ▶ **Wait-Die**: younger tx is killed when it requests a lock currently held by an older tx
- ▶ **Wound-Wait**: younger tx is killed when older tx requests a lock currently held by the younger tx
- ▶ **Wait-die** incurs more rolled back txs
  - ▶ Younger txs make relatively more lock requests than older transactions
  - ▶ But those younger txs have probably performed “less work”
- ▶ **Wound-wait** incurs fewer rolled back txs
  - ▶ They’ll be making (relatively more lock requests) and this protocol allows them to wait
  - ▶ Only aborted when older tx (which is less likely to make lock requests) tries to lock the younger tx’s data
  - ▶ But: in “wound-wait”, when we do rollback the younger tx, it will have done “more work” compared to transactions rolled back under the “wait-die” protocol
    - ▶ By definition: will be holding locks (and have done some work)




## Can “Coarse Locks” Improve Performance?

- ▶ One approach to improve lock-based performance is to increase the granularity of a lock footprint
  - ▶ Note: until now, we’ve been assuming a one-to-one mapping from database objects to locks
  - ▶ A million item  $\Rightarrow$  million locks
- ▶ What if we reduce the work that the lock manager must perform
- ▶ Specifically: model “data items” as a hierarchy of data granularities
  - ▶ Where the small granularities are nested within larger ones
  - ▶ Think of a “tree” data-structure
- ▶ Locking a “node” higher in the tree explicitly also locks lower levels of that tree implicitly
- ▶ Tradeoff:
  - ▶ Finer granularity, lower in tree  $\Rightarrow$  high concurrency but high locking overhead
  - ▶ Coarse granularity, higher in tree  $\Rightarrow$  low locking overhead but low concurrency

# DATABASE LOCK HIERARCHY



# DATABASE LOCK HIERARCHY



## A New Lock Type: Intention Locks

- ▶ The **motivation** for coarse-granularity locking is efficiency
  - ▶ Example: by locking at the **file level**, no need to lock at **record level**
- ▶ There's a major problem lurking here: can **you spot it**?
- ▶ Scenario:  $T_i$  locked  $F_b$  explicitly, now  $T_j$  wants to lock a record that's **nested in  $F_b$** 
  - ▶ How does the DBMS “know” that  $T_i$  has the lock on that record?
  - ▶ Must traverse **locking tree from root to this record** to get that information
    - ▶ This problem is not a “deal-breaker”: just an observation
- ▶ Another scenario:  $T_j$  wants to lock at a **higher level** than  $T_i$  lock granularity
  - ▶ Only way to block that request is for DBMS to traverse the **entire tree**!
- ▶ Seems to entirely negate the motivation for this idea 😊



## Intention Lock Modes Solve This Problem (I)

- ▶ **Intention lock modes** are an efficient way to record the nested locking information
- ▶ DBMS acquires **intention locks** at higher levels of the tree before acquiring an **explicit lock** at lower levels of the tree
- ▶ Example: to acquire a **row-lock**, DBMS requires that
  1. An intention lock be acquired on the **table**
  2. After the table lock has been granted, an intention lock is acquired on the **page**
  3. Only after the page lock has been granted, will the DBMS allow the tx to acquire the row lock
- ▶ Locks are acquired in **root-to-leaf** order, released in **leaf-to-root** order
- ▶ Payoff:  $T_j$  detects the intention lock without having to traverse the entire tree

## Intention Lock Modes Solve This Problem (II)

- ▶ Intention-Shared (IS) lock
  - ▶ Indicates explicit locking at a lower level with **shared locks**
- ▶ Intention-Exclusive (IX) local
  - ▶ Indicates explicit locking at lower level with **exclusive or shared locks**
- ▶ Shared + Intention-Exclusive (SIX)
  - ▶ Subtree rooted at that node is locked explicitly in **shared** mode and indicate explicit locking at a lower level with **exclusive locks**
- ▶ See textbook for details
- ▶ Also: bonus section in lecture

# COMPATIBILITY MATRIX

		$T_2$ Wants				
$T_1$ Holds		IS	IX	S	SIX	X
	IS	✓	✓	✓	✓	×
	IX	✓	✓	×	×	×
	S	✓	×	✓	×	×
	SIX	✓	×	×	×	×
	X	×	×	×	×	×

## (Intention) Lock Protocol


- ▶ Each tx obtains appropriate lock at **highest level** of the database “tree”
- ▶ To get S or IS lock on a node, the txn must hold at least IS on parent node
- ▶ To get X, IX, or SIX on a node, must hold at least IX on parent node

Intention locks allow a higher level node to be locked in S or X mode **without having to check all descendent nodes**

## Bonus: Hierarchical Locking Demo


## EXAMPLE – TWO-LEVEL HIERARCHY

Read Andy's record in R.




## EXAMPLE – TWO-LEVEL HIERARCHY

Read Andy's record in R.



## EXAMPLE – TWO-LEVEL HIERARCHY


Update Lin's record in R.





## EXAMPLE – TWO-LEVEL HIERARCHY


Update Lin's record in R.



## EXAMPLE – THREESOME


Assume three txns execute at same time:

- $T_1$  – Scan **R** and update a few tuples.
- $T_2$  – Read a single tuple in **R**.
- $T_3$  – Scan all tuples in **R**.




## EXAMPLE – THREESOME

Scan **R** and update a few tuples. **T<sub>1</sub>**




## EXAMPLE – THREESOME

Scan **R** and update a few tuples.



## EXAMPLE – THREESOME








## EXAMPLE – THREESOME

Scan all tuples in **R**.





# Today's Lecture: Wrapping it Up

Lock-Based Protocols

Deadlocks

Lock Granularities

Bonus: Hierarchical Locking Demo

- ▶ Today's lecture corresponds to the textbook **Chapter 18.1-18.3**