

# **OATSdb: V1 Milestone**

Providing Atomicity & Isolation For a Main-Memory  
Database

Avraham Leff


Yeshiva University

*avraham.leff@yu.edu*

COM3563: 2020

1. Introduction
2. Atomicity
3. Isolation
4. Testing: A Non-Definitive Discussion
5. Measuring Performance

# Introduction



In the **Vo** milestone you implemented a “**transactionally aware**” main-memory database

- ▶ A Map-based programming model
- ▶ A **server-side container** which required client to be participate in a transaction when using the programming model
- ▶ Clients can “begin”, “commit”, and “rollback” transactions
- ▶ The V1 milestone will **build upon your previous work**



In the **V1** milestone you will provide **atomicity** and **isolation** properties for your main-memory database

- ▶ These are the **A** and **I** from transactional ACID
- ▶ Specifically: you will use **locking semantics** as an implementation technique
  - ▶ The details are up to you ...
  - ▶ But: you're required to use "pessimistic" locking approach

- ▶ There is a common perception that databases are about **persistence**
- ▶ I grant that persistence is a very important feature ☺...
- ▶ But: it's independent of “atomicity” and “isolation”!
  - ▶ Sanity check: did any of the textbook or lecture algorithms ever get involved with “main-memory” *versus* “disk” issues?
- ▶ And: “main-memory” transactions definitely provide value
  - ▶ Because “all-or-nothing” semantics are very attractive
- ▶ These semantics require that the DBMS provide “atomicity” to address client-side failures or deliberate rollbacks
- ▶ These semantics require that the DBMS provide “isolation” to address multi-threaded environments

- ▶ The V1 milestone requires you to have a good understanding of **Java concurrency**
- ▶ The good news: Java has fantastic support built into the JDK
- ▶ The bad news: successive waves of increased JDK capabilities mean that there are a lot of ways to accomplish what you want
- ▶ Online information less helpful than usual: tends to assume that you already understand the big picture ☹
  - ▶ Also: such information may not have been updated even when “better alternatives” became available
- ▶ Book recommendation: **Java Concurrency in Practice**, by Brian Goetz and Tim Peierls
  - ▶ Yes, it's “old”, published in 2006
  - ▶ Still very worth reading, especially if you read the newer Javadocs as well
  - ▶ **Amazon**

Interface from Vo **will not change** despite DBMS and TxMgr functional changes

- ▶ You can & will implement as you choose
- ▶ So long as the only required public APIs are those provided by DBMS, Tx, and TxMgr
  - ▶ And you support the OATSDBType **dbmsFactory** and **txMgrFactory** methods for the V1 enum value
- ▶ See the requirements document for package and \$DIR information



# Atomicity

- ▶ Independent of **concurrency**, transactions offer the ability to “draw a box” around a chunk of code that states:
  - ▶ Do not **commit** this work until the entire chunk succeeds
  - ▶ If a failure occurs in the middle of the transaction, clients should **not see any change** from the pre-transaction state of the data
    - ▶ Semantics: as if the transaction had never executed ☺
- ▶ Clients can explicitly force the transaction to be **rolled back** through the `TxMgr.rollback()` API
  - ▶ Expectation: no code effect of that transaction will be visible to the next transaction or piece of code
  - ▶ Same semantics as unintentional failure
- ▶ Note: the DBMS may trigger a rollback for reasons of its own
  - ▶ “Stay tuned”

### Any ideas?

- ▶ Remember: there are several possible techniques

You can use any implementation technique that you want: here, the only requirement is that you provide the semantics specified above!

- ▶ What follows is only a suggestion to get you thinking
- ▶ Key idea: rollback requirement implies that the DBMS must be able to restore all client data database system to **pre-transaction state**
- ▶ At the same time, the DBMS must make client-induced data changes visible to the client as her code executes
  - ▶ Meaning: all CUD operations must behave “normally” from the client’s perspective
  - ▶ (And as we provide **isolation**, those CUD operations should be “invisible” to **other clients**)
- ▶ How can we accomplish these seemingly contradictory goals?

- ▶ Not my terminology (see [Wikipedia](#))
  - ▶ Not clear whether my usage is broadly adopted ...

- ▶ Key idea: apply client's **ongoing** tx work to a **copy of the database**
- ▶ I'll refer to this "copy" as a **shadow db**
  - ▶ If tx commits, propagate shadow db state to "real" db
  - ▶ If tx rolls back, "no harm, no foul", just silently dispose of the shadow db 😊

Issue you'll want to consider:

- ▶ Do we need to provide clients with a complete copy of the database?
- ▶ Do we need to provide clients with a copy of data that they've only **read** (not accessed via a CUD operation)?

- ▶ This approach really shines with respect to handling tx failure and rollback 😊
- ▶ But: how do we propagate the client's work to the “real” database when a transaction **commits**?
- ▶ Conceptually: want to (atomically) **replay** all of the tx's work against the “real” db
- ▶ What would such a log look like?
- ▶ At a minimum: need to log the method **name** and the method **parameters**
- ▶ But: that's not enough!
  - ▶ Remember: we're using a Java programming model
  - ▶ Method state involves more than the method parameters!
  - ▶ Method “state” also involves the **entire** state of the object instance on which the method was invoked
- ▶ Can **you suggest** how we can propagate the shadow database to the real database?

- ▶ (Almost) all Java objects can be **serialized**
  - ▶ Meaning: converted into a **stream of bytes**
- ▶ Take a look at the **Java Object Serialization Specification** (or a textbook)
- ▶ Serialized objects can be subsequently **deserialized**
  - ▶ Meaning: the stream of bytes are converted (back) to a live object instance
- ▶ So an object can be “flash frozen” now and “reconstituted” later
- ▶ Example scenarios:
  - ▶ Save the bytes to a file
  - ▶ Transmit the bytes across the network
  - ▶ And: **for the V1 milestone** ...
    - ▶ Copy the bytes from the “real db” → “shadow db”
    - ▶ Copy the bytes back from the “shadow db” → “real db”
  - ▶ Bytes are bytes ☺

- ▶ You don't need anything subtle from Java serialization in this project
  - ▶ You're not allowing external clients to use serialization to pass you objects with **possibly risky data**
    - ▶ That can be a serious **security problem**
- ▶ And: "we'll cheat" and not worry about ensuring that **object evolution (versioning)** doesn't break your code
- ▶ Plenty of material on the 'net on how to do the basics ...
- ▶ (see my previous link to the serialization protocol ...)
- ▶ Make sure you understand Java's approach to **Object Streams**

Key point: we're using serialization as an easy (for us) way to make **object copies on demand**

**Values** must be serializable because they are “copied via serialization” from real Map to shadow Map (and back again)

- ▶ **Q:** do Map **Keys** have to be serializable as well?
  - ▶ Why or why not?
  - ▶ **A:** for V1: **Keys do not need to be serializable**
  - ▶ Because they stay constant during the transaction!
  - ▶ Only **Values** associated with keys may change
- ▶ **Note:** this description (misleadingly) makes serialization sound like work

*A Java object is serializable if its class or any of its super-classes implements either the `java.io.Serializable` interface or its subinterface, `java.io.Externalizable`.*

### Java Tutorial

- ▶ You'll find that most objects are serializable “out of the box”



- ▶ The **OATSdb** programming model requires that database resources only be accessed within a transaction
- ▶ **Q:** Can we prevent the client from breaking the rules with the following (trivial) scenario?

```
TxMgr.begin();  
// get reference to Map, then get reference to Account instance  
final Account account = AccountMap.put("key", 42);  
TxMgr.commit();  
// Client accesses Account instance (outside a tx!)  
final int value = account.get("key");
```

- ▶ **A:** the “shadow db” approach also protects the DBMS from this scenario
- ▶ Serialization ensure that your clients only get access to “shadow Map” and “shadow Object” references
- ▶ The “Real map” and “real Object” reference **are different references**!
- ▶ These semantics follow from the fact that “serialization” creates brand new references
  - ▶ State of the new reference is identical to original object
  - ▶ But the references are different!

### Important caveat:

- ▶ Unfortunately, we can't **prevent** the client from using that ill-gotten Account reference 😞
- ▶ Ideally, we'd like to provide semantics that permit clients to
  1. Acquire a reference in  $tx_1$
  2. Use that reference in  $tx_2$
- ▶ ...With the DBMS silently “swizzling” the **client's reference** to point to the “real” **server's reference**
- ▶ Similar semantics to “single-level store” semantics: client's shouldn't have to worry about transaction subtleties!
- ▶ Note: you **must provide** precisely these semantics for Map references!
  - ▶ Map references acquired in a committed (or rolled back)  $tx_1$  are valid in  $tx_2$

- ▶ **Q:** if we want this “single-level” store semantics ...
  - ▶ And we are providing these semantics for Map references ...
  - ▶ Why not insist on these semantics for Map.Entry references as well?
- ▶ **A:** because this would require a lot of work on your part ☺
- ▶ In **Vo**, you must have figured out a way to provide clients with a “Map” reference that’s really a reference to an **enhanced Map** object
  - ▶ One that’s transactionally aware in a way that a “vanilla” Map reference is unable to be
- ▶ That’s fairly easy to do because the Map API is “known in advance”
- ▶ Requiring similar behavior for **arbitrary Map.Entry** references requires a “code generation” step that must precede deploying code to your DBMS

- ▶ After consultation with experts, I decided that a “code generation” phase requires too much work for too little pedagogical insight ☹
- ▶ Also: the straightforward implementation technique would still constrain the programming model to require that DBMS objects have **interfaces**
  - ▶ That constraint is unacceptable: e.g., would disallow String and Long values
- ▶ Bottom line: clients must acquire fresh references to Map.Entry instances on a per-transaction basis
  - ▶ The DBMS **doesn't have to detect** that a client is violating this rule
  - ▶ But: any work performed by a client on a “stale” reference is wasted!
    - ▶ State modifications performed on the stale reference are not performed on either the “real” or a “shadow” copy

If you can think of a way to solve these issues neatly, please speak to me (after class) ☺

## Atomicity: Summarizing The Requirements

- ▶ You absolutely **do not** have to adopt the implementation approach we've been discussing!
- ▶ The **isolation requirements** are the following:
  - ▶ Work performed in  $tx_1$  is only visible to  $tx_2$  **after  $tx_1$  commits**
  - ▶ If  $tx_1$  fails or is rolled back, its work **will not be visible to  $tx_2$**
- ▶ All of the VO requirements apply “as is” to V1, modulo the reference semantics discussed above
  - ▶ Map references do not have to be (re)acquired in a new tx
  - ▶ Map.Entry references do have to be (re)acquired in a new tx
    - ▶ Otherwise:  $tx_1$  work performed on a stale reference will not be visible to  $tx_2$
  - ▶ But: the DBMS **is not responsible** for throwing an exception if a client uses a stale Map.Entry reference
- ▶ Note: these atomicity requirements are **independent of isolation**!
  - ▶ Here  $tx_2$  refers to a transaction that executes **after**  $tx_1$  completes

- ▶ Clients never get a reference to a “real db object”
  - ▶ Only to a “shadow db object”
- ▶ On a **per-tx basis**, DBMS grows a set of “shadow db” objects that are associated with this transaction instance
  - ▶ One set of objects **per Map instance**
- ▶ All of the client’s work is performed on the shadow database!
- ▶ The DBMS transparently creates the “shadow db” objects using Java Serialization on behalf of the client
- ▶ `Tx.commit()`: the DBMS copies the shadow database state back to “real db”, **using Serialization** to protect the “real db” from the client’s current set of references
- ▶ `Tx.rollback()`: the DBMS discards “shadow db” objects so that they are not visible to the “real db”

# Isolation

- ▶ **Atomicity** is about changing database state in “all or nothing” fashion
- ▶ **Isolation** is about protecting one database user from another
- ▶ In our “single user” implementation, this may seem artificial
  - ▶ But the algorithms are the same as for the “multi-user” implementation
- ▶ The problem is all about **concurrency**: intuitively, “how does the system protect two users from trashing each other’s work”



- ▶ **Pessimistic**: assume that users **will** concurrently access the same data
  - ▶ DBMS must protect one from the other in advance
  - ▶ As discussed in lecture, pessimistic approach is implemented with a family of **lock-based** protocols
- ▶ **Optimistic**: assume that users **will not** concurrently access the same data
  - ▶ DBMS verifies that assumption when transaction commits
  - ▶ Rolls transaction back if assumption turns out to be false
- ▶ In V1, you must implement a **lock-based protocol** to provide **isolation**
- ▶ I urge you to read a good book on the concurrency facilities provided by the JDK
  - ▶ Followed by multiple readings of the [java.util.concurrent Javadocs](#) and [java.util.concurrent.locks Javadocs](#)
- ▶ Otherwise: you likely will reinvent the wheel, badly ☹

- ▶ Assume that the DBMS “locks” a resource on behalf of  $tx_1$  such that  $tx_2$  cannot access that resource
- ▶ What is the result of  $tx_2$  invoking “get”, “put”, etc on the locked resource?
  - ▶ **Q**: should the DBMS throw an Exception?
  - ▶ **A**: absolutely not!
  - ▶ Think of all the  $tx_2$  work that will be wasted simply because  $tx_1$  is (temporarily) holding that lock
  - ▶ Instead:  $tx_2$  **blocks**, waiting for DBMS to **unlock** the resource
- ▶ **Q**: when should the DBMS unlock the resource?
- ▶ **A**: when  $tx_1$  either commits or rolls back ...

- ▶ Naive approach: lock the **entire database** on a per-tx basis
  - ▶ Advantages: easy to implement 😊
  - ▶ Disadvantages: resulting performance will be atrocious ☹
- ▶ Slightly less naive: lock on **per-Map** granularities
  - ▶ Similar advantages and disadvantages as above
- ▶ Can we do better?
- ▶ Sure: at a minimum, lock on **per MapEntry** granularities
  - ▶ Motivation: this will increase DBMS concurrency because  $tx_1$  and  $tx_2$  can execute **concurrently**
  - ▶ Note: the finer granularity will not provide benefit if the two transactions access the same resource concurrently

You must implement locking at MapEntry granularities!

- ▶ As you know from lecture, we can further refine a lock-based protocol by allowing resources to be locked in one of two modes:
  - ▶ **Shared mode:**  $tx_1$  can only read the resource
    - ▶ And  $tx_2$  can **also read** that resource
    - ▶ Potentially increasing concurrency considerably
  - ▶ **Exclusive mode:**  $tx_1$  can both read & write the resource
    - ▶ No other  $tx_2$  can even read that resource
    - ▶ Reduces concurrency, but required to provide transactional **isolation**

**OATSdb V1 simplification:** all locks are exclusive locks

### Motivation:

- ▶ Given `Map.get` semantics, would be very difficult for the DBMS to detect when a lock must be “upgraded”
  - ▶ There is no API through which the client declares that she’s about to **modify the state** of an object
  - ▶ She can acquire a `Map.Entry` reference via `Map.get`, then modify the reference through **that object’s API**
- ▶ Good news: less work for you 😊
- ▶ But be aware of the performance implications: **OATSdb** will be unable to take advantage of applications with high read/write ratios 😞

- ▶ V1 will implement the strict 2PL locking protocol (refer back to lecture)
- ▶ Advantages: guarantees **transactional serializability** without requiring that you generate and analyze transaction schedules
  - ▶ Note: “tx serializability” has totally different meaning from the Java “serializability” we’ve been referring to 😊
- ▶ Implication: if you use the 2PL protocol, you’re guaranteeing transactional isolation

- ▶ DBMS acquires lock on transactional resource “on demand” for  $tx_i$ 
  - ▶ Blocks if another  $tx_j$  has previously acquired a lock on that resource
  - ▶ Proceeds only after lock is released
- ▶ Transactional commit or rollback releases all locks acquired in  $tx_i$
- ▶ Where should you implement locks?
- ▶ Suggestion: at the boundary between the “real db” and the “shadow db”!
  - ▶ Atomicity implementation: client operates only on “shadow db” resource
  - ▶ Isolation implementation: client allowed to copy from “real db” to “shadow db” only if client first acquires a lock on that resource

- ▶ Because V1 is using a pessimistic, lock-based, protocol, client transactions can potentially **deadlock**
- ▶ Scenario:
  - ▶  $tx_1$  access  $MapEntry_i$  and  $MapEntry_j$  **in that order**
  - ▶  $tx_2$  access  $MapEntry_j$  and  $MapEntry_i$  **in that order**
  - ▶ DBMS grants locks in the order that the transactions request them
  - ▶ Both transactions get “stuck”: no further progress possible ☹️
- ▶ From lecture you know that basically, two approaches possible
  - ▶ **Prevent** deadlocks
  - ▶ **Detect** and recover from deadlocks
  - ▶ Each of these high-level strategies has many, many, lower-level implementations



**OATSdb** approach: Implement “deadlock prevention” by using a **timeout-based** scheme

1.  $tx_1$  requests a lock at time  $t_1$ , but is blocked by the DBMS
2. DBMS starts a clock ticking with a system-specified **timeout period**
  - ▶ Timeout expires at  $t_1 + \text{SystemTimeout}$
3. If timeout expires and  $tx_1$  still hasn't acquired the lock it needs ...
  - ▶ The DBMS **rolls back  $tx_1$** , and releases all locks currently held by  $tx_1$
  - ▶ (This allows **other transactions** to make progress)
4. Else ...
  - 4.1 The transaction previously holding the lock **releases that lock**
  - 4.2 The DBMS “wakes up”  $tx_1$ , allowing it to acquire the lock and resume execution

```
public interface ConfigurableDBMS extends DBMS {  
    /** Sets the duration of the "transaction timeout".  
     * A client whose transaction's duration exceeds  
     * the DBMS's timeout will be automatically rolled  
     * back by the DBMS.  
     *  
     * @param ms the timeout duration in ms, must be  
     * greater than 0  
     */  
    void setTxTimeoutInMillis(int ms);  
  
    /** Returns the current DBMS transaction timeout  
     * duration.  
     *  
     * @return duration in milliseconds  
     */  
    int getTxTimeoutInMillis();  
}
```

Your V1 DBMS must implement this interface

- ▶ Key problem: what's the optimal timeout duration?
- ▶ Set it **too high**  $\Rightarrow$  limits concurrency
- ▶ Set it **too low**?  $\Rightarrow$  may result in very irritated (rolled back) clients
  - ▶ Their transactions may have actually been making progress, just taking a long time
- ▶ Keep these serious issues in mind, but ...
- ▶ Don't worry about them for your V1 implementation!
- ▶ I'll simply supply your DBMS with an arbitrary timeout value, you'll plug that value into your lock-based concurrency control implementation

- ▶ Clients must **acquire a lock** on each MapEntry that her transaction accesses
  - ▶ Whether via “get”, “put”, “remove”
- ▶ Transaction execution **blocks** if a lock cannot be acquired
  - ▶ Presumably because another transaction has acquired the lock on that resource and not yet completed
- ▶ Transactions “queue” for resource locks
  - ▶ DBMS must ensure **fairness**
- ▶ Happy path: tx acquires locks on all its resources and commits
- ▶ Otherwise: tx **times out** and DBMS rolls back the transaction
  - ▶ This scenario manifests as a  
`edu.yu.oatsdb.base.ClientTxRolledBackException` to the client

## Scenario

1. Tx acquires locks on all resources
2. Tx enters a long running computation
  - ▶ Or an infinite loop 😊
3. Or throws an exception
4. What happens to transactions blocked on this transaction's resource set?
  - ▶ They're doomed 😞
  - ▶ DBMS has no way to rollback another thread that's executing "arbitrary code"
  - ▶ Programmers must understand this and adopt a "release locks as fast as possible" approach
    - ▶ Burden is on the transaction code
  - ▶ Alternatively: DBMS can run a periodic sweep at **transaction granularity** to detect whether a transaction is taking too long
    - ▶ Release locks as necessary & rollback the transaction
    - ▶ **Not responsible for this feature in V1**

The “block while acquiring lock” approach relies on the fact that Java threads:

- ▶ Can sleep (suspend their own execution)
- ▶ And receive notifications that they were interrupted (told by another thread to do something else)
- ▶ These Java thread properties are the basis of any solution for “block & resume” implementation!
- ▶ This is true whether or not you directly exploit these properties or whether you use JDK classes that exploit these properties ...
- ▶ But: please, please, don't reinvent the wheel 😊

- ▶ Tx thread sleeps the system-specified duration
- ▶ When sleeper awakes, DBMS makes *n* attempts to acquire the lock on behalf of that transaction
- ▶ After *n* unsuccessful attempts, rolls back the transaction
- ▶ Advantage: simplicity
  - ▶ See `Thread.sleep` API
  - ▶ But: consider the serious problems with this approach ☹
- ▶ Disadvantage: performance (throughput)
  1.  $tx_1$  acquires lock
  2.  $tx_2$  blocks and sleeps for ten minutes
  3.  $tx_1$  commits **two seconds after  $tx_2$  goes to sleep**
- ▶ This naive strategy doesn't provide a way for the DBMS to proactively **wake up** the sleeping thread
  - ▶ Result: throughput takes a hit ☹

- ▶ DBMS **aggressively wakes up** sleeping transaction
- ▶ Whenever  $tx_1$  either commits or rolls back ...
  - ▶ DBMS selects another  $tx_2$  (if any) that is waiting for a  $tx_1$  resource
  - ▶ Wakes up  $tx_2$
  - ▶  $tx_2$  now acquires that lock and resumes execution
- ▶ Advantage: improves throughput because tx will block for the  **$\min(\text{timeout}, \text{lock release})$** 
  - ▶ Timeouts are an **upper-bound** on duration of tx blocking
  - ▶ Also could be used for “zombie transactions” (see above, not required for V1)
- ▶ Disadvantage: implementation can't simply do a “sleep with a timeout”



# Testing: A Non-Definitive Discussion

- ▶ Can **create** Maps
  - ▶ But not if **previously created**
  - ▶ But not if name is **already associated with a different Map**
- ▶ Can **retrieve** those Maps
  - ▶ But only if **previously created**
- ▶ Can insert a MapEntry
- ▶ Can retrieve that MapEntry
- ▶ Can delete that MapEntry

All of the above **only if client thread is currently associated with a transaction!**

## Is Map Type Safety Enforced?

- ▶ Can I trick your DBMS into returning a Map with different **key** class?
- ▶ Can I trick your DBMS into returning a Map with different **value** class?
- ▶ Can I trick your DBMS into returning a Map with **different name**
  - ▶ With the same key & values classes as previously created Map ...
- ▶ Does your DBMS enforce the requirement that value class must be **serializable**
- ▶ Can I trick your DBMS into allowing client to supply an “empty” (null, or only whitespace) name?

- ▶ All access to TxMgr API must be done from **inside a transactional scope**
  - ▶ Only exception is: begin ☺
  - ▶ And: getTx, getStatus
- ▶ Once transaction commits, client is no longer in a transaction
  - ▶ Unless she starts a new transaction
- ▶ Ditto if the client does a transaction **rollback**
- ▶ Does your DBMS prevent **nested transactions**?
- ▶ Have you correctly implemented the required **transaction state transitions**?

- ▶ Is “put then get” activity visible **after tx commits**?
  - ▶ Meaning: visible to another transaction
  - ▶ (Important: all of these scenarios invoke the Map methods in **different – but sequential – transactions**)
- ▶ What about a “put then get then modify” scenario
  - ▶ Verifying that once committed, that object state becomes part of the “real db” ...
- ▶ What about a “put then remove then get” scenario
  - ▶ That “get” should return `null`
- ▶ What about a “put then remove & put then get” scenario
  - ▶ The “put” (taking place in the same transaction as the “remove”) should supersede the effect of the “remove”
- ▶ Is a reference to a MapEntry rendered “useless” after tx commits?
  - ▶ Meaning: can’t “cheat” by modifying that object **outside transactional scope**

- ▶ Is “put then get” activity **not visible** after tx is **rolled back**?
- ▶ Is an reference to a MapEntry rendered “useless” after tx is rolled back?
  - ▶ Meaning: can’t “cheat” by modifying that object **outside transactional scope**

The above suggestions represent a set of requirements that I **may or may not choose to validate** when I test your code.

Furthermore, the above suggestions in no way preclude the use of **different tests** of your code.

- ▶ These above suggestions are essentially “JUnit” tests
  - ▶ They are “single-threaded” tests
  - ▶ The single thread **serially** simulates the behavior of **serial transactions**
- ▶ They apply to the **atomicity** part of the V1 milestone

### Definition

A **multi-threaded** test is a test that involves multiple, **concurrently executing** client threads.

**Q:** Do you need to perform multi-threaded tests against **OATSdb**?

**A: Absolutely!** The point of this exercise is to have a transactional system, one that provides atomicity and isolation to clients **in the face of concurrent activity** from database clients



- ▶ It is **very hard** to use JUnit (or similar test harnesses) to do **multi-threaded** tests
- ▶ There are various approaches out there
  - ▶ My reaction: much more trouble than they're worth ☹
  - ▶ You're more than welcome to use any system you choose

- ▶ Please **let me know** if you find something that “works”
- ▶ Meaning:
  - ▶ A library that lets me focus on my test semantics
  - ▶ A library that includes “tests setup” APIs that are easy to understand and are unobtrusive

- ▶ You want to set up a number of concurrent clients
- ▶ Have them stress-test your system
  - ▶ Issuing “get/remove/put” requests
  - ▶ Against the same data
- ▶ Your test code needs to detect whether the **OATSdb** **system under test** failed to provide its transactional guarantees
  - ▶ Example: did one thread modify a MapEntry concurrently with another thread's actions?
  - ▶ Example: did a transaction involving multiple MapEntries violate atomicity?

I already have that test, and it's not hard for you to create your own

## Some classes that I found helpful ...

```
import java.util.concurrent.Callable;  
import java.util.concurrent.ExecutionException;  
import java.util.concurrent.ExecutorService;  
import java.util.concurrent.Executors;  
import java.util.concurrent.Future;  
import java.util.concurrent.FutureTask;  
import java.util.concurrent.TimeUnit;  
import java.util.concurrent.atomic.AtomicLong;
```

1. Create some “example Map resources”
2. Create a sample transaction that uses these resources
3. Instrument them to report a concurrency violation
4. Have a thread-pool of clients that spin up concurrent transactions
5. Step back and let the sparks fly

- ▶ Advantage of “stress test” approach: either your DBMS buckles under the pressure or it doesn’t ...
- ▶ Disadvantage: what do you do if (when?) your DBMS fails
  - ▶ It’s very, very difficult to debug concurrent code
  - ▶ And: your test results are **non-deterministic**
    - ▶ May pass, then fail, then pass
    - ▶ One failure trace will differ from the next
- ▶ You’ll be tempted to simply ignore the problem
  - ▶ Especially if it’s intermittent
  - ▶ Don’t: this will cost you points
  - ▶ And bad job appraisals in your career ☹
- ▶ Approach: augment the stress test with a set of **precision, multi-threaded** tests
  - ▶ Each of which focuses on a specific scenario

We want each test to:

- ▶ Be a **reproducible** unit test for that scenario
- ▶ That implies fairly strict control of a few threads
- ▶ Pare back the scenario to its essentials
- ▶ Then sprinkle **lots** of print statements throughout the code
  - ▶ But not so many that they interfere with the concurrency
- ▶ Key idea: specify the concepts that you **can assert** about the correct behavior of your DBMS
  - ▶ Assert them! Never rely on “println validation” 😊
- ▶ Let's walk through one of my examples ...
  - ▶ Again: you **don't have to do it this way**
  - ▶ You may well come up with a better approach!

Let's call it: **Tx2BlocksWhileTx1Locks**

- ▶ Your DBMS must block  $tx_2$  from accessing a resource that is currently locked by  $tx_1$ 
  - ▶ Note: stress-test may not detect a failure because  $tx_1$  may access, then release the lock, so quickly that difficult to detect that  $tx_2$  was ever blocked in the first place
- ▶ There are different things that can go wrong with this scenario
- ▶ Here we verify that
  - ▶  $tx_2$  is in fact blocked by  $tx_1$  lock
  - ▶  $tx_2$  will **will not be rolled back** if  $tx_1$  releases the lock before the timeout period elapses
  - ▶ (Separate test that verifies that  $tx_2$  **is rolled back** if  $tx_1$  holds on to that lock for too long)

## Implementing Two Transactions

- ▶ You **don't need a special** transactional resource and a transactional “program” for this sort of test
  - ▶ Such as you do need for “stress test” code
- ▶ All you need is minimal code that invokes `begin` and `commit`
  - ▶ You're not testing the **state** of the resource, just the locking behavior (**sequencing**)
- ▶ You can specify the timeout for your DBMS under test
  - ▶ Example: normally might be a minute, you can set it to 500ms
- ▶ Code up an (inner) class that will implement  $tx_1$  behavior and another that will implement  $tx_2$  behavior
  - ▶ Have each implement `Runnable`
  - ▶ Now you can specify (with **reasonable milli-second accuracy**) how long the associated tx thread should **sleep**

- ▶ Create the Maps and transactional resource(s) that  $tx_1$  and  $tx_2$  will access
  - ▶ Will have to do this in a transaction ☺
- ▶ Create a data-structure into which the transactions can record the time that they perform various actions of interest
- ▶ Your test will involve **assertions** about the **order** in which the transactions create these records
- ▶ Your test will involve **assertions** about the “content” of these records
  - ▶ Example: what's the timestamp of an individual record
- ▶ You definitely **don't** want this data-structure to be a transactional resource ☺
- ▶ Advice: use **static inner classes** to keep your code in one Java class, and to allow sharing of the global instance variables



1. Begin tx
2. Access transactional resource
3. Sleep for **half a timeout period**
  - ▶ Key point:  $tx_1$  is blocking  $tx_2$  from accessing this resource
  - ▶ Even though  $tx_1$  isn't "using" the resource, all resources accessed by this transaction are locked until  $tx_1$  commits or aborts
4. When thread wakes up, commit the tx
  - ▶ Key point: this unlocks the resources that were previously locked by the transaction
  - ▶ We expect  $tx_2$  to be able to access the resource

Have the thread create (and record) appropriate data that allows you to re-create the event timeline

1. Sleep a bit at the beginning to ensure no “race condition” with  $tx_1$
2. Begin tx that involves accessing the same resource that  $tx_1$  is using
3. Q: how long should  $tx_2$  block for a v1 implementation?
4. When  $tx_2$  is permitted to access the resource, record the time
  - ▶ We can assert that
    - 4.1 That  $tx_2$  will access the resource
    - 4.2 And: assert (approximately) what time it did so
5. Commit the transaction

Have the thread create (and record) appropriate data that allows you to re-create the event time-line

- ▶ Have your `main` configure the timeout period you want
- ▶ Create `tx1` and `tx2` instances
- ▶ Create Thread instances that wrap these “transaction” instances
  - ▶ Tip: you can `setName` to assign useful names to the threads you’re creating
- ▶ Start the threads
- ▶ When the threads finish executing, perform your “assert” logic and determine whether the test passed or failed

Q: Can you see a problem here?

- ▶ `main` will start and finish **before** `tx1` and `tx2` have **finished executing** 😞
- ▶ Any suggestions?
- ▶ We could have the `main` thread sleep for our “best guess” as to how long the transaction threads will take
  - ▶ Plus a bit more for safety
- ▶ Better approach: use a `java.util.concurrent.CountDownLatch`
  1. Initialize to number of transaction threads
  2. `main` invokes `latch.await`
    - ▶ Will not proceed with program execution until latch value reaches 0
  3. Each transaction thread invokes `latch.countDown` (just before) it finishes execution
  4. `main` will “automagically” resume processing

# Measuring Performance

- ▶ We've focused so far on **transactional correctness**
  - ▶ Specifically: does the DBMS provide **atomicity** and **isolation** guarantees
- ▶ Suggestion: (fairly easy) implementation
  - ▶ Allow exactly one client to use the DBMS at a time
    - ▶ Note: this will not provide “atomicity”, just “isolation”
- ▶ **Q:** what's wrong with the above approach?
- ▶ **A:** we **also want our implementation** to have good performance 😊
  - ▶ Which we'll define here in terms of throughput
    - ▶ Meaning: “number of committed transactions per second”

- ▶ Not difficult to do the basics ...
  - ▶ Create transactional resources
  - ▶ Create transactional program
  - ▶ Wind them up, and measure throughput ...
- ▶ We need at least one “measurement knob”: must specify the test’s **concurrency factor**
  - ▶ Defined as the *“number of transactions executing at the same time”*
  - ▶ How do you expect performance to change as you turn this “knob” one way or the other?
- ▶ Also must think about:
  - ▶ How many transactional resources will you create?
  - ▶ Does it even matter how many you create?

## Scenario:

- ▶ I create a 100,000 resources
- ▶ My transactional program accesses only a single resource
- ▶ I create 10 threads, each of which **randomly** selects a single resource to access in its transaction
- ▶ There is **almost no thread contention for DBMS resources** at all!
  - ▶ Resulting numbers are “useless”
  - ▶ I can increase the number of threads and get increasingly better numbers without paying a penalty for quite a while
- ▶ We need to devise a mechanism that can serve as a “tunable knob” on our infrastructure



- ▶ What if we set the number of transactional resources to some constant?
  - ▶ Almost any value will do
- ▶ Intuitively: by varying the **number of resources** used by each transaction, we can increase or decrease the **amount of competition** between transactions
  - ▶ That number can be expressed as a **percentage** (of the **constant** number of resources available)
- ▶ **Tx Footprint Ratio**:  $accountsPerTx/nAccounts$ 
  - ▶ The larger the ratio, the more inter-transactional competition
  - ▶ And this ratio is **independent** of the actual number of transactional resources!

- ▶ We definitely want to measure throughput
- ▶ Should report the total number of transactions **leaving** the system
  - ▶ Be useful if we start to wonder whether we're running the measurement long enough
- ▶ Should report the number of "successful" transactions
  - ▶ Defined as "txs that committed"
- ▶ And the number of "failed" transactions
  - ▶ Defined as "txs that were rolled back"
  - ▶ Key point: our transactional logic will not **willingly** rollback a tx
  - ▶ So this number reflects the number of txs rolled back by the system because of "supposed" **deadlocks**

- ▶ Think about this issue: what is the role of timeout duration in your performance infrastructure?
  - ▶ This factor seems to be important as it interacts with the **average duration** of transaction execution
  - ▶ Transaction duration is itself (partly) a function of the **number of transactional resources per transaction**
  - ▶ Do we (and how would we) make these tunable knobs?
- ▶ Also: I urge you to test your **intuition** against **observed phenomena**
- ▶ Example: given observed performance for a **concurrency factor** of 1 ...
  - ▶ How do you predict performance to change as we **increase the concurrency factor**?
  - ▶ How do you predict performance to change as we **increase the transaction footprint ratio**?
- ▶ If your intuition doesn't match your observations
  - ▶ Is this a “teachable moment”?
  - ▶ Or: a bug in your code?

- ▶ Getting your system “correct” is most important
- ▶ Getting good performance is less important (for now) but still very important
  - ▶ For one thing: otherwise you can “cheat” by reducing the opportunities for concurrency so as to decrease the likelihood of isolation failures ☺

Note: I am seriously considering assigning **bonus project points** based on how your implementation performs **relative** to e.g., average of class performance

Introduction

Atomicity

Isolation

Testing: A Non-Definitive Discussion

Measuring Performance